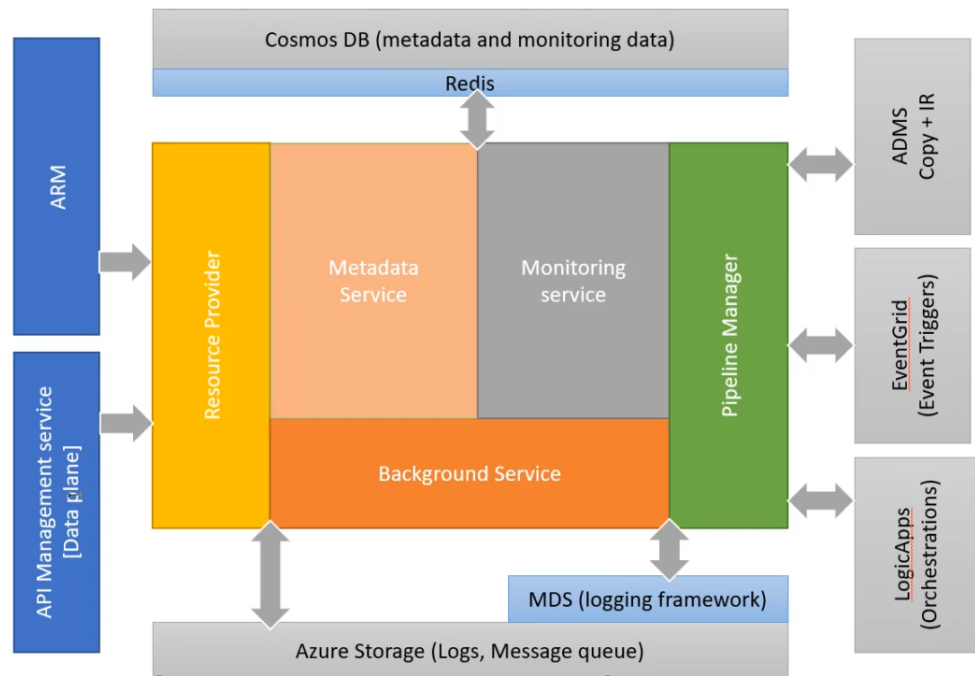


ADMS Kusto Queries and Performance Kusto Functions

Last updated by | Veena Pachauri | Mar 8, 2023 at 11:57 PM PST

ADMS Kusto Queries and Performance
Kusto Functions
Thursday, November 28, 2019
10:23 AM



- ARM: HttpIncomingRequests, HttpOutgoingRequests, ARMPR_V2
- ADF RP, MS, PM, MON: ApiOperationEvent
- ADF RP, MS, PM, MON: AdfTraceEvent
- ADMS: ExecutionApiCall
- ADMS: CustomLogEvent

Normal Kusto query

```
ExecutionApiCall
| where PipelineJobId contains "edf8c715-7872-4909-bc72-26541ccb3f"
```

ExecutionApiCall can be used to check all the activity based on the Pipeline job id.

After that, you can get the activity ID with CustomLogEvent to get more detail about the activity status.

Below table is used to check reported event log on the IR Server which included the callstack to help us investigate further.

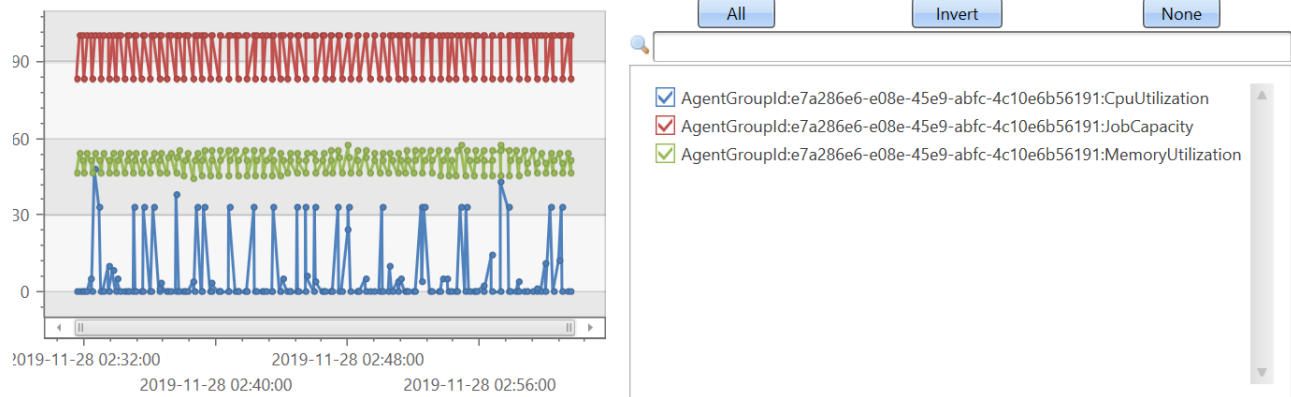
```
TraceGatewayLocalEventLog | where UserReportId contains "dcb92348-d948-4882-b534-90a7e01d98c5"
```

Below Function is used to check perf related issue.

```
202da45b-84d2-4827-8cd0-a8a18cb566d0 is the activity ID.
```

```
DiagnosticsSHIRPerfCountersById('@'202da45b-84d2-4827-8cd0-a8a18cb566d0')
```

DiagnosticsSHIRPerfCountersById kusto query is used to query the IR capacity, if the capacity is quite high than expected such as selfHost-IR, we need to scale up the self-IR node. If jobcapacity is very l is good, scale up by increasing the number of concurrent jobs that a node can run. You might also want to scale up when activities time out because the self-hosted IR is overloaded. As shown in the follow: increase the maximum capacity for a node:



Integration Runtime: IR-CONTOSO-IT				
Settings	Nodes	Auto update	Sharing	
Name	Status	IP Address	Limit Concurrent Jobs	Actions
Node_2	Running	Get IP Address	5	
Node_1	Running	Get IP Address	10 <input checked="" type="checkbox"/>	

From <https://icm.ad.msf.net/imp/v3/incidents/details/160817085/home>

<https://github.com/MicrosoftDocs/azure-docs/blob/master/articles/data-factory/create-self-hosted-integration-runtime.md>

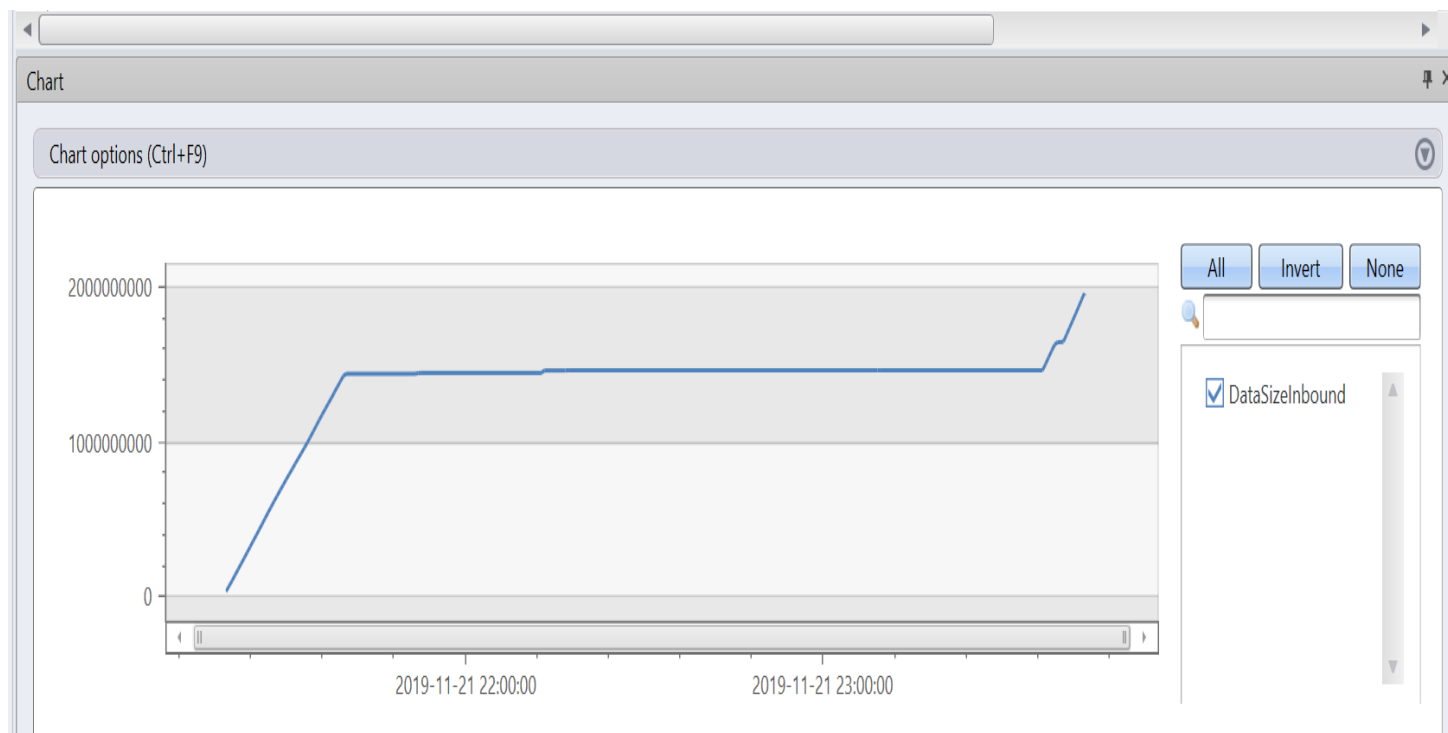
ShowCopyTimelineByActivityId

ShowCopyTimelineByActivityId is used to check the Job status for each stage.

TIMESTAMP	JobCreated	JobPickedUp	CopyInQueue	CopyStarted	CopyEnded	
2019-11-21 21:19:32.7457037			1			
2019-11-21 21:19:33.0575854	1					
2019-11-21 21:19:37.6842752		1				
2019-11-21 21:19:37.7390000				1		
2019-11-21 23:44:19.7390000					1	
2019-11-21 23:44:20.1575223	1					
2019-11-21 23:44:26.3520084		1				
2019-11-21 23:46:42.1699716	1					
2019-11-21 23:46:43.2623130		1				

DiagnosticsCopyThroughputById(@'202da45b-xxxx-xxxx-xxxx-xxxxxxxx')

DiagnosticsCopyThroughputById is used to get the copy status throughout, you can see the line without any change which means something stuck there from source as Inbound status. Then we can get dump and driver log to investigate further.



As you can see below as example that the throughput is bigger at beginning since the inbound growing much within minutes with enough buffer, that's why you can see throughput is bigger at beginning, after a while, due to outbound growing little, so caused the inbound reached to stable as well not growing too much after buffer full.

The throughput cx see on portal = $\text{DataSizeInbound} / \text{duration}$. That's why the value keeps going down.

Setting of batch size as 50000

General
Source
Sink
Mapping
Settings
User properties

Sink dataset *
knoffercon
Open
New

Write behavior
Insert

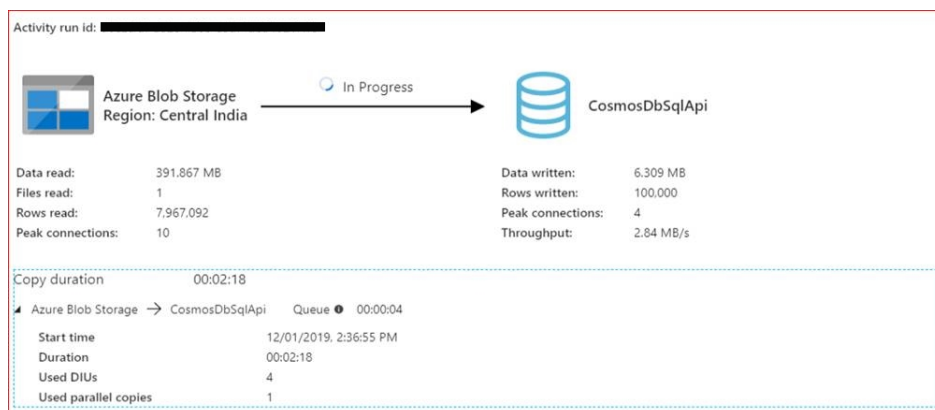
Write batch timeout

Write batch size
50000

Max concurrent connections

Disable performance metrics analytics
☐


Copy started with 3 MB/s



Gradually reduced to 1.5 MB/s


[Learn more on copy performance details from here.](#)

Activity run id:



Azure Blob Storage
Region: Central India

In Progress



CosmosDbSqlApi

Data read:

Files read:

Rows read:

Peak connections:

424.243 MB

1

8,625,298

10

Data written:

Rows written:

Peak connections:

Throughput:

25.239 MB

400,000

4

1.443 MB/s

Copy duration 00:04:54

▲ Azure Blob Storage → CosmosDbSqlApi Queue ● 00:00:04

Start time

Duration

Used DIUs

Used parallel copies

12/01/2019, 2:36:55 PM

00:04:54


4

1

Details Refresh


[Learn more on copy performance details from here.](#)

Activity run id:



Azure Blob Storage
Region: Central India

In Progress



CosmosDbSqlApi

Data read:

Files read:

Rows read:

Peak connections:

436.54 MB

1

8,875,287

10

Data written:

Rows written:

Peak connections:

Throughput:

41.013 MB

650,000

4

978.155 KB/s

Copy duration 00:07:37

▲ Azure Blob Storage → CosmosDbSqlApi Queue ● 00:00:04

Start time

Duration

Used DIUs

Used parallel copies

12/01/2019, 2:36:55 PM


00:07:37

4

1


Decreasing less to 500 KB/s

Activity run id:



Azure Blob Storage
Region: Central India

In Progress



CosmosDbSqlApi

Data read:

Files read:

Rows read:

Peak connections:

468.51 MB

1

9,525,257

10

Data written:

Rows written:

Peak connections:

Throughput:

82.024 MB

1,300,000

4

587.934 KB/s

Copy duration 00:13:36

▲ Azure Blob Storage → CosmosDbSqlApi Queue ● 00:00:04

Start time

Duration

Used DIUs

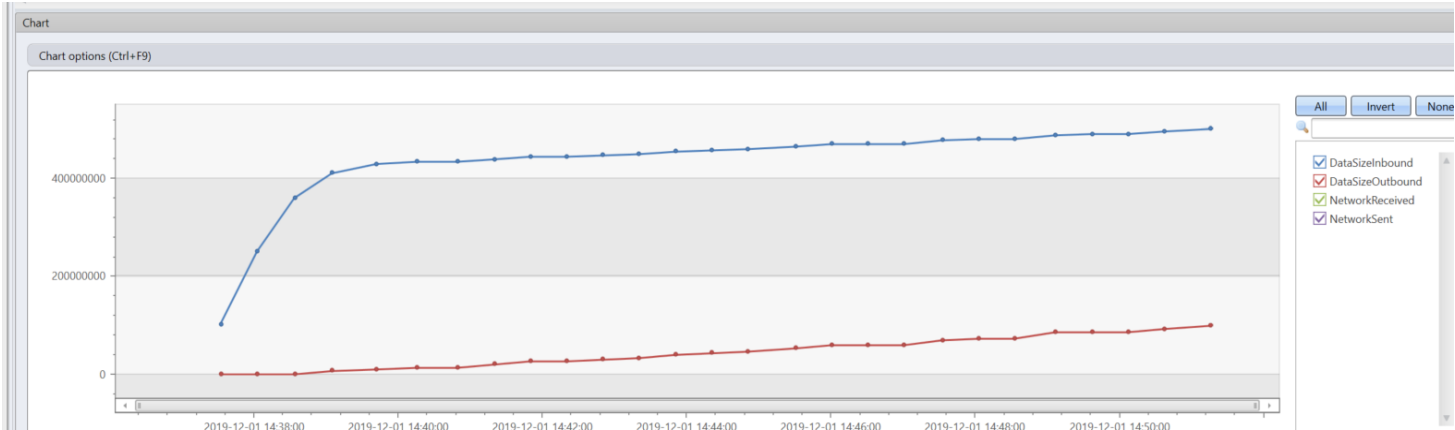
Used parallel copies

12/01/2019, 2:36:55 PM

00:13:36

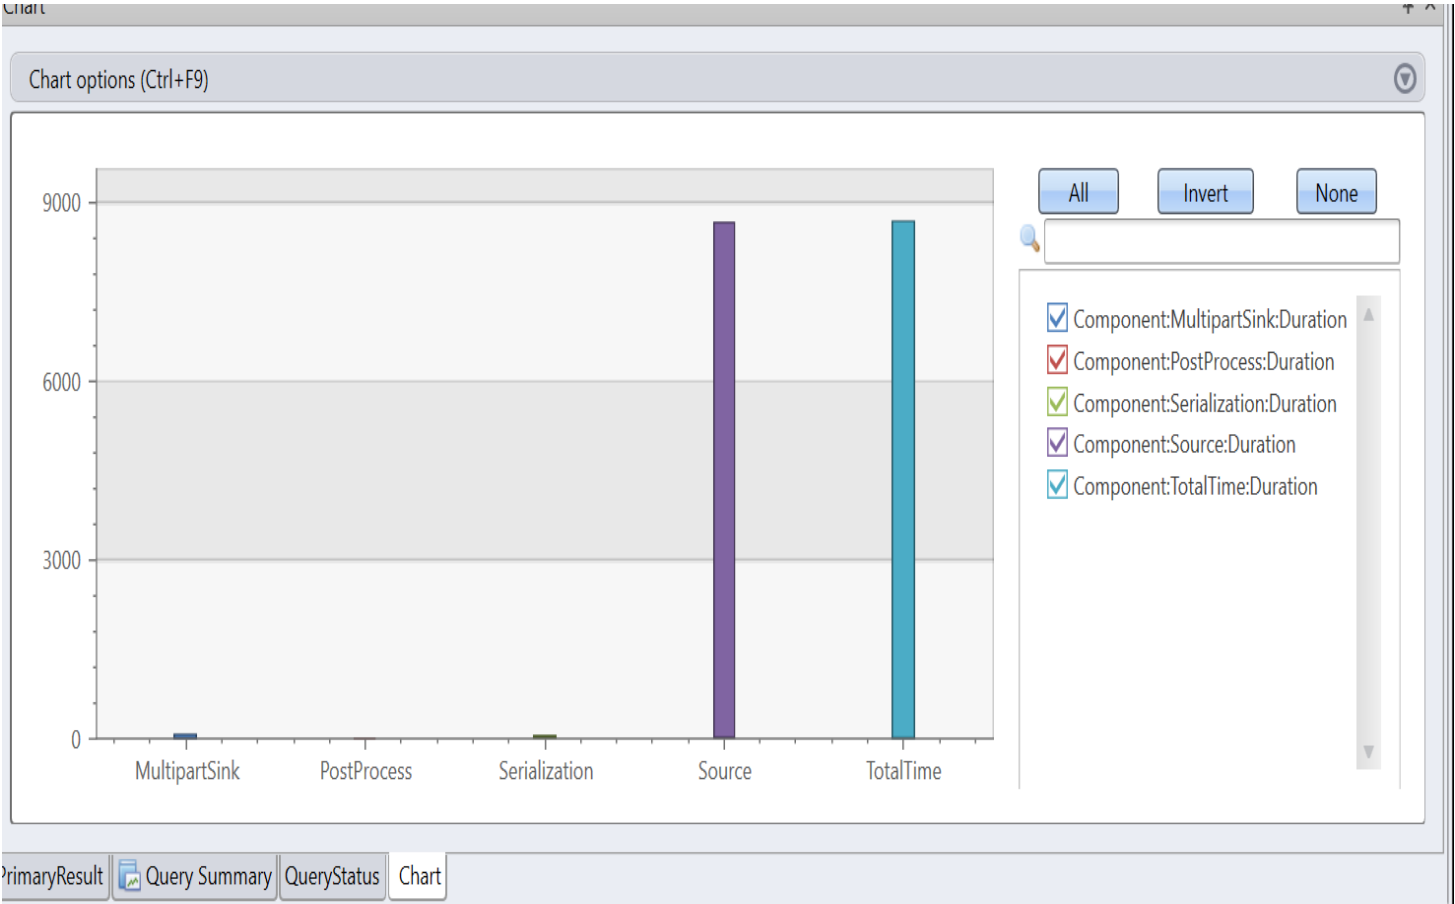
4

1



ShowCopyPerfBottleneckByActivityId(@'202da45b-xxx-xxxx-xxxx-xxxxxxx')

ShowCopyPerfBottleneckByActivityId is used to check where the bottleneck is, from below chart, you can see it was stuck from Source side, so you need to figure out any issue between IR and Source



Please also refer to the following TSG to further investigate perf issue during copy activity.

[\[Diagnostic\].\[V2\].Identify performance bottleneck of Copy Activity.run](#)

Note: We have merged the kusto query ShowCopyTimeLineByActivityId, DiagnosticsSHIRPerfCountersById, DiagnosticsCopyThroughputById and ShowCopyPerfBottleneckByActivityId to ASC, it will support the Chart like below:

Performance: Copy Timeline

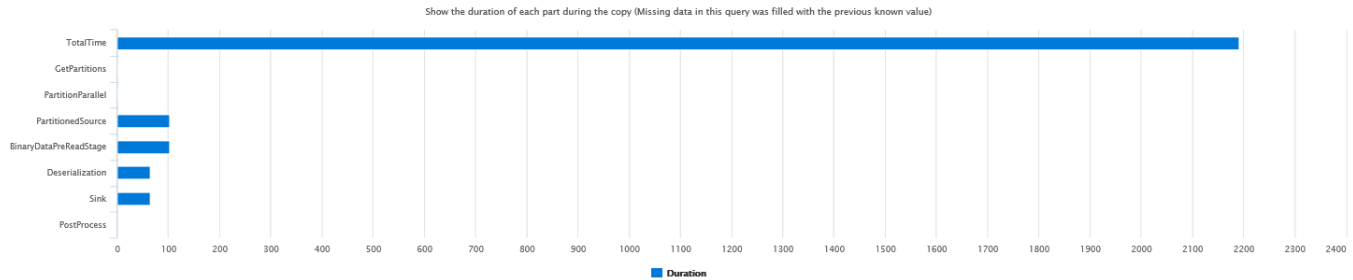
Show checkpoints during copy activity (eg: queued, job picked up, job finished)

Drag a column header and drop it here to group by that column

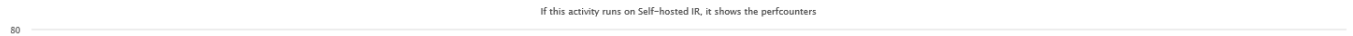
TIMESTAMP	JobCreated	JobPickedUp	CopyInQueue	CopyStarted	CopyEnded
> 2019-12-02 23:12:45				1	
> 2019-12-02 23:12:45			1		
> 2019-12-02 23:12:46	1				
> 2019-12-02 23:12:46		1			
> 2019-12-02 23:49:15					1

1 - 5 of 5 items

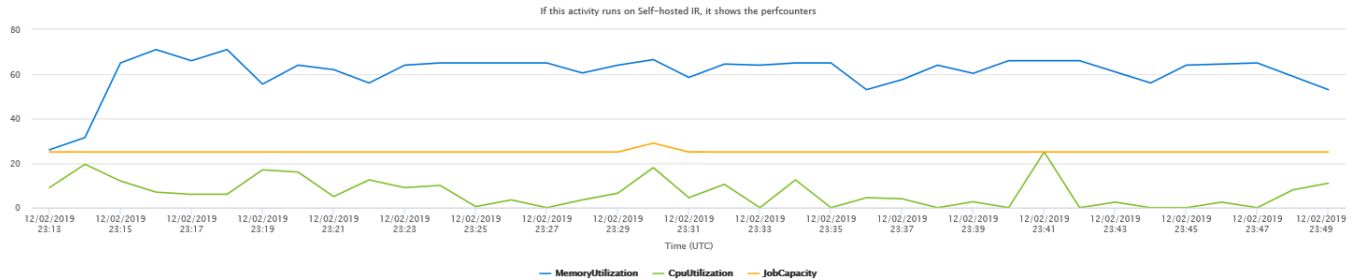
Performance: Copy Bottleneck



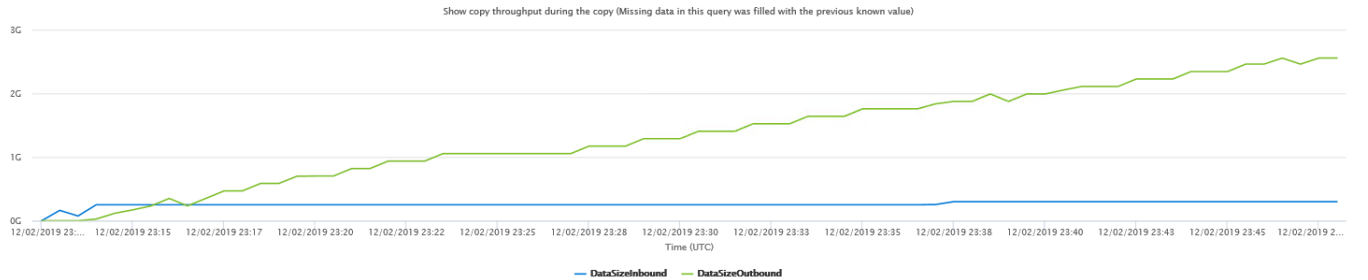
Performance: SH IR Capacity



Performance: SH IR Capacity



Performance: Copy Throughput



Created with Microsoft OneNote 2016.

How good have you found this content?

