

Data Preview of blob container failing from Storage Event Trigger

Last updated by | Anil K B | Feb 16, 2023 at 11:41 PM PST

Contents

- [Symptoms](#)
- [Cause](#)
- [RCA](#)
- [Resolution](#)
- [Trigger Configurations](#)
- [Other Prerequisites:](#)

Symptoms

- Customer is trying to create a storage event trigger to a storage account with **Selected networks in the storage account settings**. They are seeing the below error when trying to access the container inside the storage account using **Storage Event trigger**.

Data preview



Make sure you have specific filters. Configuring filters that are too broad can match a large number of files created/deleted and may significantly impact your cost.

Event Trigger Filters

Container name: **input**

Starts with:

Ends with:

0 blobs matched in "input"

Refresh

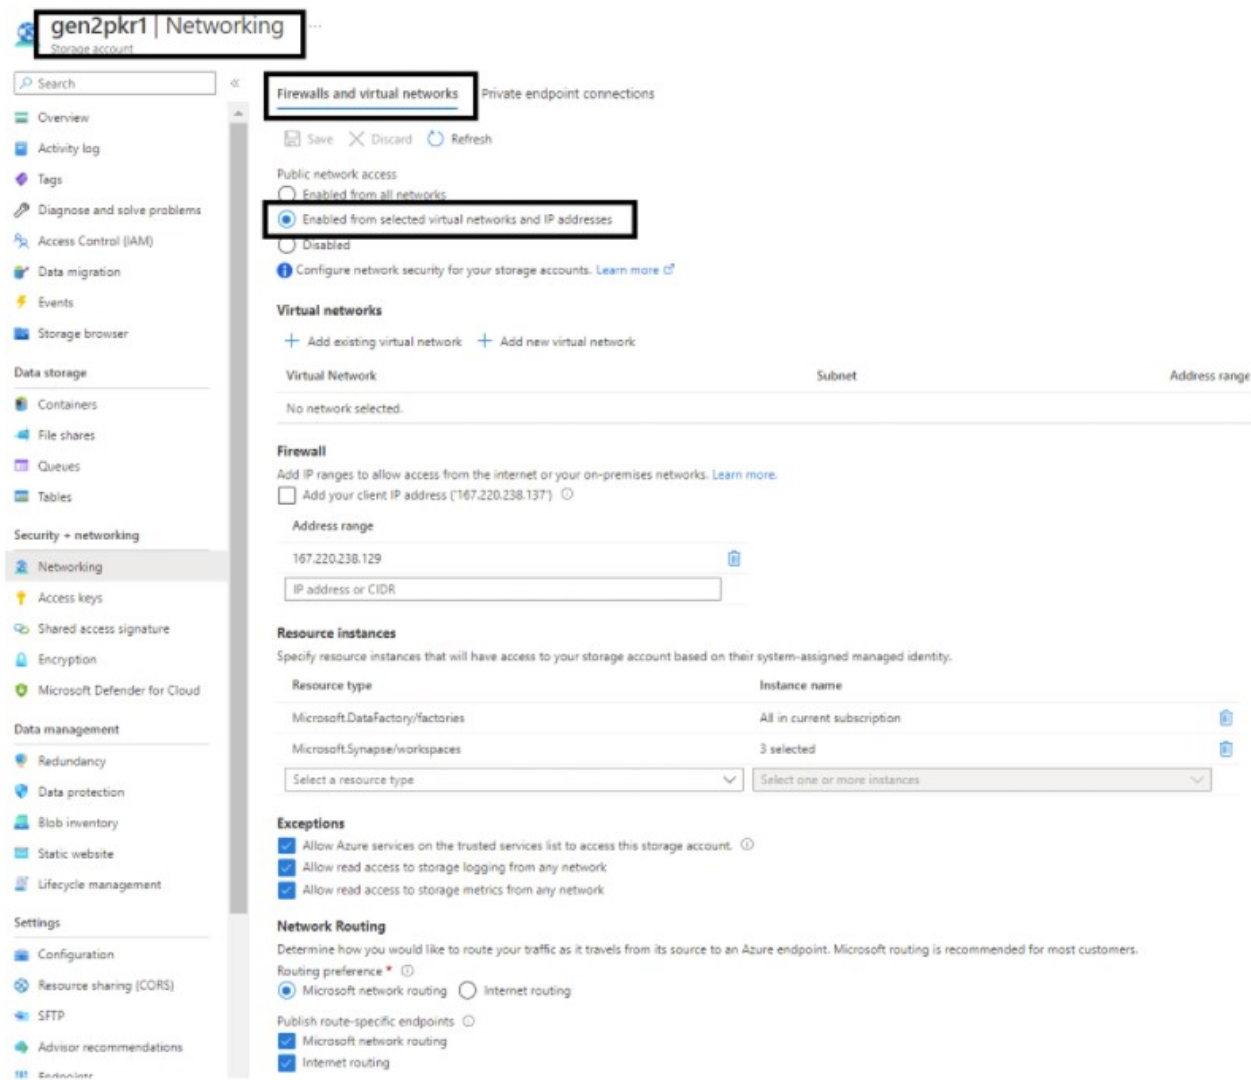
Unable to list files and show data preview. Please ensure that:



- * Your storage account is not using a firewall
- * You have the 'Storage Blob Data Reader' role, or, permission to list the storage access keys.

Cause

- The storage account is in “**Enabled from selected virtual network and IP Addresses**”

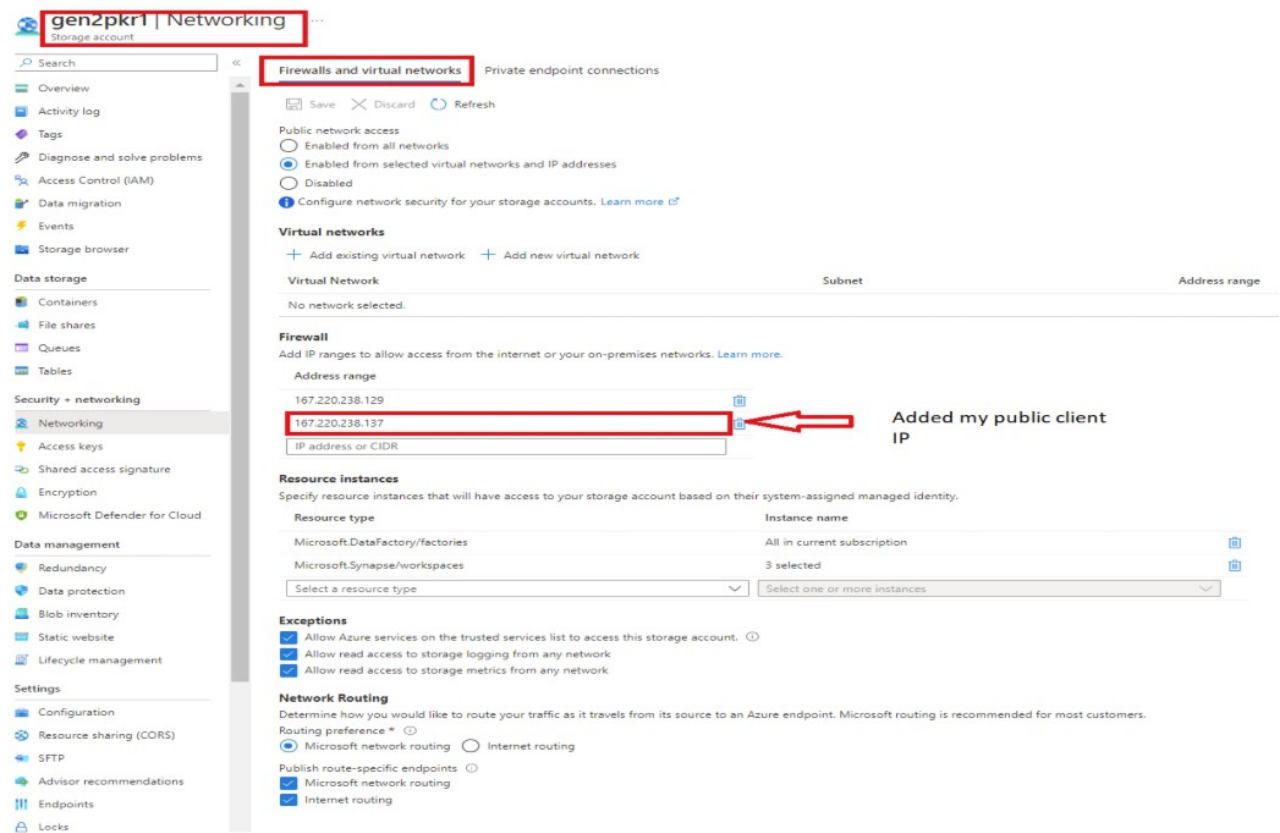


RCA

- The key point is we are enabling access of storage from selected VNET and IP addresses the requests which originate from these IP address would be the ones which would be successful when it comes to any operation on the storage account .
- Now if we don't whitelist one of those IP addresses (in this case machine where ADF is run from where the request originates) the ADF won't be able to access the blob. We can use a different client (without any whitelist) and try to run a LIST blob API and we would hit the same error there as well.

Resolution

- If the ADF is placed in **Public Endpoint** the below steps to be followed
 - Machine IP address changes every time we change a network hence we need to **whitelist the new IP address** every time we want to do data preview from storage account.



- Will be able to list the blobs in Storage Events Data Preview

Data preview

Make sure you have specific filters. Configuring filters that are too broad can match a large number of files created/deleted and may significantly impact your cost.

Event Trigger Filters

Container name: **input**

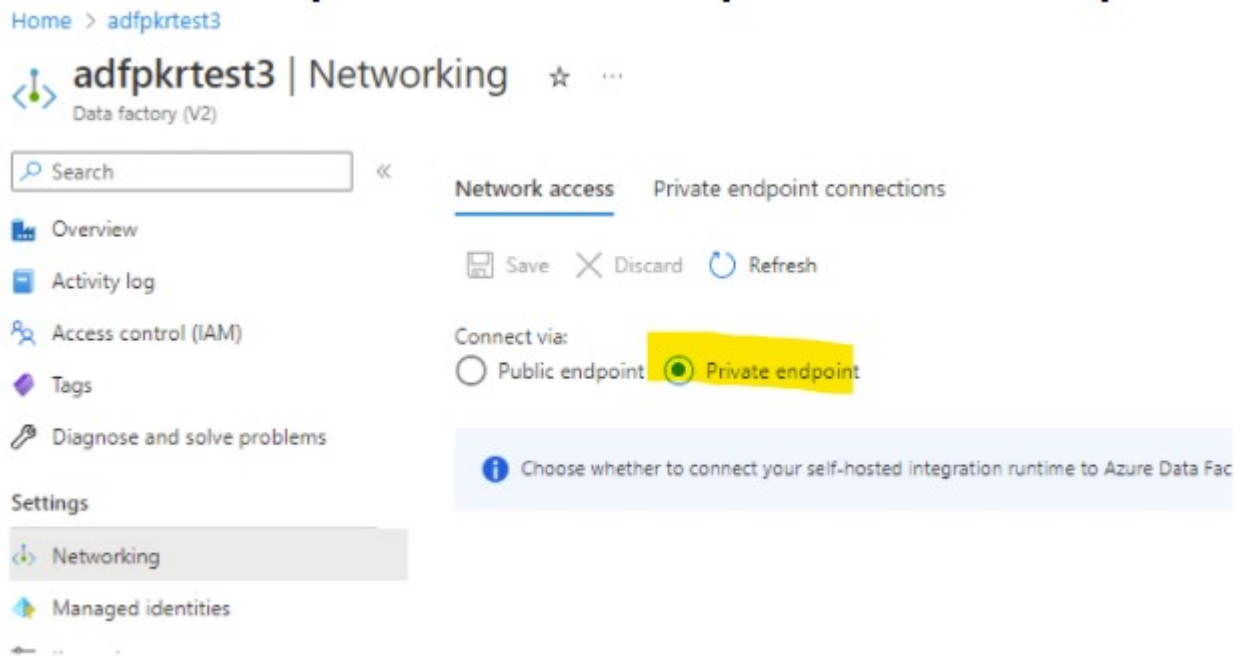
Starts with:

Ends with:

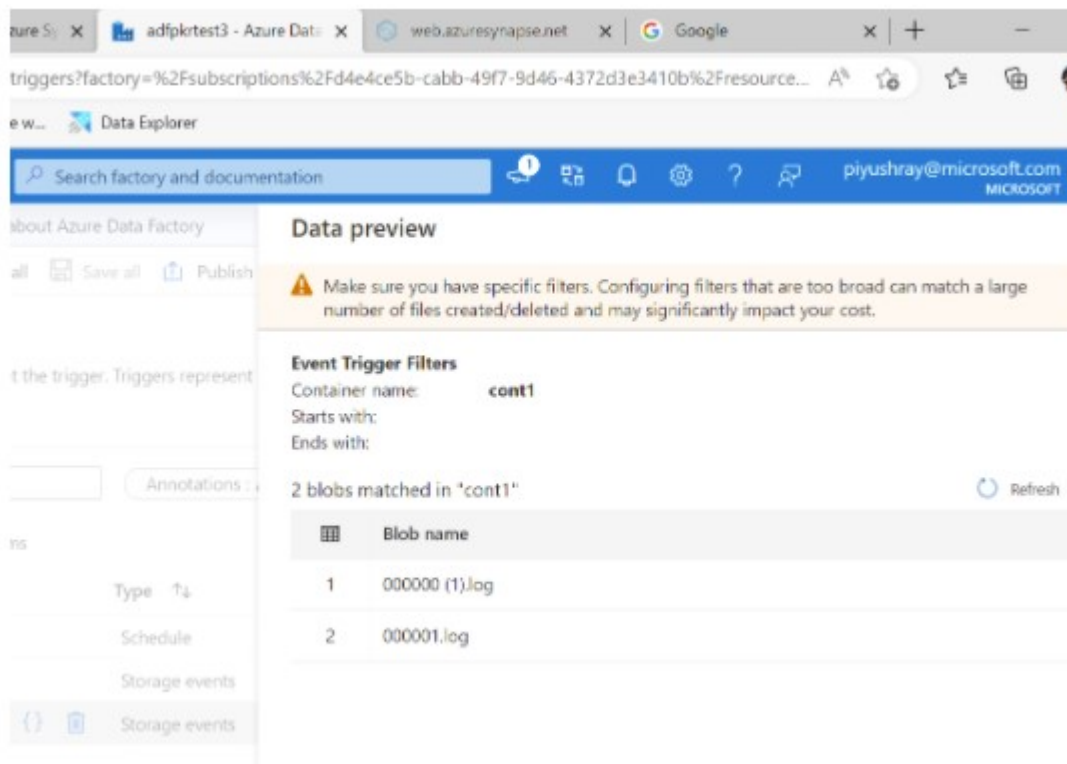
11 blobs matched in "input" Refresh

	Blob name
1	2017.csv
2	EANPool1.xlsx
3	Events.xml
4	SF_LFA_Invoices_Extract.csv
5	View_SF_LFA_Invoices.xlsx
6	start.txt
7	synapse
8	synapse/workspaces
9	synapse/workspaces/pkrsynapsews
10	test1
11	test1/meetingAttendanceReport(ADF fails on weekends - TrackingID#2204070030001463).csv

2. If the ADF is placed in **Private Endpoint** the below steps to be followed



- You need to ensure you have created portal and dataFactory private endpoint in the same VNET and you have placed another VM machine on the same VNET.
- Then go to the Networking settings on your storage account and whitelist the VNET which contains the ADF Private Endpoints(portal and dataFactory) and the Virtual Machine on which you are accessing the ADF privately.
- Once done take a RDP to the VM which is on the same VNET and the open ADF portal.
- Now if you try to list the blobs you can successfully list them over the private setup.





Trigger Configurations

- Define the trigger and set the "**Blob Path begins with**" or "**Blob Path ends with**" in accordance with the way the blob is going to be uploaded.

The screenshot shows the 'Trigger Configuration' window for a 'BlobEventsTrigger'. The 'Name' field is 'trigger2'. The 'Description' field is empty. The 'Type' is 'BlobEventsTrigger'. Under 'Account selection method', 'From Azure subscription' is selected. The 'Azure subscription' is 'PiyushRay_Azure Internal Subscription (d4e4ce5b-cabb-49f7-9d46-4372d3e3410b)'. The 'Storage account name' is 'gen2plr1'. The 'Container name' is 'input'. The 'Blob path begins with' field is 'start', which is highlighted with a red rectangular box. The 'Blob path ends with' field is empty. Under 'Event', 'Blob created' is checked. 'Ignore empty blobs' is set to 'Yes'. There is a '+ New' button under 'Annotations'. The 'Status' is 'Started'.

- In this case we have used "start" as my "Blob path begins with" and saved the definition.
- When we dropped a file named "start" in the container "input" it triggered the pipeline and the pipeline run started.

Other Prerequisites:

- Ensure **Microsoft.EventGrid** resource provider is registered in the subscription. Here is the [link](#) .
- Ensure the user has **Microsoft.EventGrid/EventSubscriptions/Write permission**. Here is the [link](#)  for reference

Additional Information

- Icm References:
- cases: 2301190040003843
- Author: **Piyush Ray**
- Reviewer: **kbanil**

How good have you found this content?

