

DataFlow Activities Time To Live (TTL) and IR configurations scenarios and recommendations

Last updated by | Ranjith Katukojwala | Jan 31, 2023 at 4:01 PM PST

Contents

- [Scenario: 1](#)
- [Scenario: 2](#)
- [Scenario: 3](#)
- [How to check Cluster Compute Acquisition Duration:](#)
- [How TTL Works?](#)
 - [Possible Questions and Recommendations:](#)
- [Additional Information:](#)

By default, every data flow activity spins up a new cluster based on the IR configuration. Cold cluster start-up time takes a few minutes (~2-3 mins) and data processing can't start until it is complete. If most of the data flow executes in parallel, it is not recommended to enable TTL.


Please note that the Warm cluster takes around ~5-10 sec to pick up the next job.

The customer could set up Dataflow IR TTL in the following Scenarios:

Scenario: 1

Customers could have a pipeline with multiple Dataflow activities with an Integration Runtime (IR) with a Time To Live (TTL) of 0 minutes or DefaultIntegrationRuntime. In this case, a separate cluster will be created for each Dataflow activity and start-up time takes a few minutes for each cluster + processing time.

For instance, the below Pipeline consists of three Dataflow activities and each activity set up with DefaultIntegrationRuntime. In this case, each DF activity will have its own Cluster.



Activity runs

Pipeline run ID **XXXXXXXX-XXXX-XXXX-XXXX-XXXXXXXXXXXX**

All status ▾					
Showing 1 - 3 of 3 items					
Activity name	Activity type	Run start ↑↓	Duration	Status	Integration runtime
Advworksdf	ExecuteDataFlow	9/10/20, 1:48:35 PM	00:04:53	✔ Succeeded	DefaultIntegrationRuntime (West US)
dataflow3	ExecuteDataFlow	9/10/20, 1:43:26 PM	00:05:09	✔ Succeeded	DefaultIntegrationRuntime (West US)
cosmostoparquet_limitedvalues	ExecuteDataFlow	9/10/20, 1:38:33 PM	00:04:53	✔ Succeeded	DefaultIntegrationRuntime (West US)

To check further details:

```
let pipelinerunid = 'xxxxxxxx-xxxx-xxxx-xxxx-xxxxxxxxxxxx';
cluster('adfcus.kusto.windows.net').database('AzureDataFactory').ActivityRuns
| union cluster('adfneu.kusto.windows.net').database('AzureDataFactory').ActivityRuns
| where pipelineRunId == pipelinerunid
| where activityType == 'ExecuteDataFlow'
| summarize activitystart = min(start), activityend = max(end),ActivityCompletion = max(end)-min(start), Activ
```

Below result would give the time taken to complete each activity (ActivityCompletion = Cluster Startup time + Processing time)

activityRunId	activityName	activityType	activitystart	activityend	ActivityCompletion	ActivityCompletionInMilliseconds
XXXXXXXX-XXXX-XXXX-XXXX-XXXXXXXXXXXX	cosmostoparquet_limitedvalues	ExecuteDataFlow	2020-09-10 20:38:33.3984138	2020-09-10 20:43:25.6890539	00:04:52.2906401	292291
XXXXXXXX-XXXX-XXXX-XXXX-XXXXXXXXXXXX	dataflow3	ExecuteDataFlow	2020-09-10 20:43:26.4911929	2020-09-10 20:48:34.6479080	00:05:08.1567151	308156
XXXXXXXX-XXXX-XXXX-XXXX-XXXXXXXXXXXX	Advworksdf	ExecuteDataFlow	2020-09-10 20:48:35.6277987	2020-09-10 20:53:28.1983484	00:04:52.5705497	292571

```
let pipelinerunid = 'xxxxxxxx-xxxx-xxxx-xxxx-xxxxxxxxxxxx';
cluster('adfcus.kusto.windows.net').database('AzureDataFactory').ActivityRuns
| union cluster('adfneu.kusto.windows.net').database('AzureDataFactory').ActivityRuns
| where pipelineRunId == pipelinerunid
| where activityType == 'ExecuteDataFlow'
| where status != 'Queued' and status != 'InProgress'
| project activityRunId,effectiveIntegrationRuntime,managedVNetName,activityName,pipelineRunId
|join kind= rightouter (
cluster('adfcus.kusto.windows.net').database('AzureDataFactory').DataflowClusterLogs
| union cluster('adfneu.kusto.windows.net').database('AzureDataFactory').DataflowClusterLogs
| where TraceCorrelationId == pipelinerunid
| where Caller == 'com.microsoft.datafactory.dataflow.AdmsClient.processedEvent'
| distinct ActivityRunId, Message,TraceCorrelationId
| extend EventProcessedTime = format_datetime(unixtime_milliseconds_todatetime(toint(extract("Processing Time
| project ActivityRunId,activityName,effectiveIntegrationRuntime,managedVNetName, Message,EventProcessedTime
```

Below result would give the Integration Runtime (IR) used and Time is taken to complete each event in the activities. (EventProcessedTime = Time taken for that event to process).

Activity Run Id	activity Name	effective Integration Runtime	managed VNet Name	message	Event Processed Time
		DefaultIntegrationRuntime (West US)		AdmsClient.processedMetric for JobId: metric: Target = identificationAttributes Processing Time = 9391 ms Event Timestamp = 2020-09-10 20:43:04.169	00:00:09.3910000
		DefaultIntegrationRuntime (West US)		AdmsClient.processedMetric for JobId: metric: Target = baseUOM Processing Time = 9391 ms Event Timestamp = 2020-09-10 20:43:04.169	00:00:09.3910000
		DefaultIntegrationRuntime (West US)		AdmsClient.processedMetric for JobId: metric: Target = fulfilmentProductState Processing Time = 9391 ms Event Timestamp = 2020-09-10 20:43:04.171	00:00:09.3910000
		DefaultIntegrationRuntime (West US)		AdmsClient.processedMetric for JobId: metric: Target = Parquetdest Processing Time = 12338 ms Event Timestamp = 2020-09-10 20:43:04.676	00:00:12.3380000
		DefaultIntegrationRuntime (West US)		AdmsClient.processedMetric for JobId: metric: Target = source1 Processing Time = 3014 ms Event Timestamp = 2020-09-10 20:53:08.867	00:00:03.0140000
		DefaultIntegrationRuntime (West US)		AdmsClient.processedMetric for JobId: metric: Target = sink1 Processing Time = 18058 ms Event Timestamp = 2020-09-10 20:53:09.109	00:00:18.0580000
		DefaultIntegrationRuntime (West US)		AdmsClient.processedMetric for JobId: metric: Target = source1 Processing Time = 3102 ms Event Timestamp = 2020-09-10 20:48:12.548	00:00:03.1020000
		DefaultIntegrationRuntime (West US)		AdmsClient.processedMetric for JobId: metric: Target = sink1 Processing Time = 19772 ms Event Timestamp = 2020-09-10 20:48:12.849	00:00:19.7720000

Here Cluster Startup time + Event Processed Time = ~ ActivityCompletion. Please check below Activity screenshot for the reference:

cosmostoparquet_limitedvalues

Data flow

Refresh

Auto refresh

On

identificationAttribu...

Sink: ●

fulfilmentProductSt...

Sink: ●

baseUOM

Sink: ●

identificationAttribu...

Sink: ●

Union

Sink: ●

Parquetdest

✓

Cluster startup time: 3m 55s 398ms

Number of transforms: 5

Parquetdest

Sink

Total columns

5

New columns

0

Updated columns

5

Dropped columns

0

Drifted columns

0

Stream information

Rows calculated

3

Total partition

3

Stage time

9s 391ms

Last update (PDT)

9/10/2020, 1:43:04 PM

Partition chart

Row count

1.0

0.9

0.8

0.7

0.6

0.5

0.4

0.3

0.2

0.1

0

1

2

3

Partition

Skewness

-

Kurtosis

-

Sink processing time

12s 338ms

Parquetdest

✓

COLUMN

METHOD

ORIGINAL SOURCE

fulfilmentProductState

Mapped

fulfilmentProductState(fulfilmentProductState)

identificationAttributes

Mapped

identificationAttributes(identificationAttributes)

imageUrl

Mapped

baseUOM(imageUrl)

baseUOM

Mapped

baseUOM(baseUOM)

packagingType

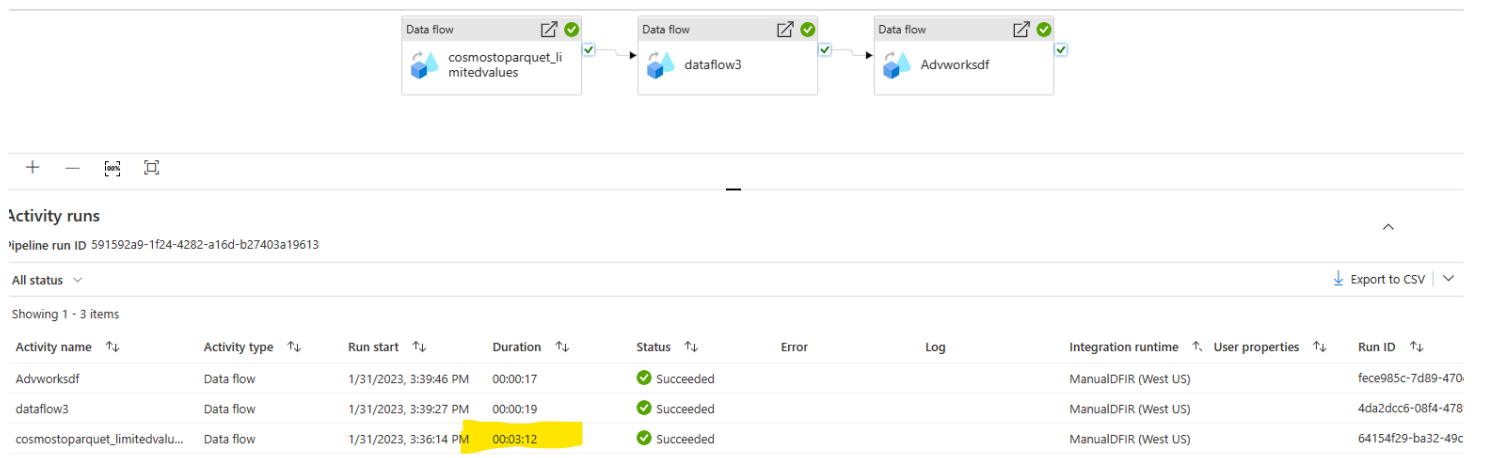
Mapped

baseUOM(packagingType)

Scenario: 2

Customers could have a pipeline with multiple Dataflow activities with an Integration Runtime (IR) with specified Time To Live (TTL). If the pipelines contain multiple sequential data flow, if we enable a time to live (TTL) value, specifying a time to live value keeps a cluster alive for a certain period of time after its execution completes. If a new job starts using the IR during the TTL time, it will reuse the existing cluster, and start-up time will greatly reduce. After the second job completes, the cluster will again stay alive for the TTL time. Only one job can run on a single cluster at a time. If there is an available cluster, but two data flows start, only one will use the live cluster. The second job will spin up its own isolated cluster.

For instance, the below Pipeline consists of three Dataflow activities with each activity setup with an Azure Integration Runtime with TTL of 10 Minutes. **Note:** The recommendation is to use the smallest TTL, more TTL minutes incur more cost.



Activity runs

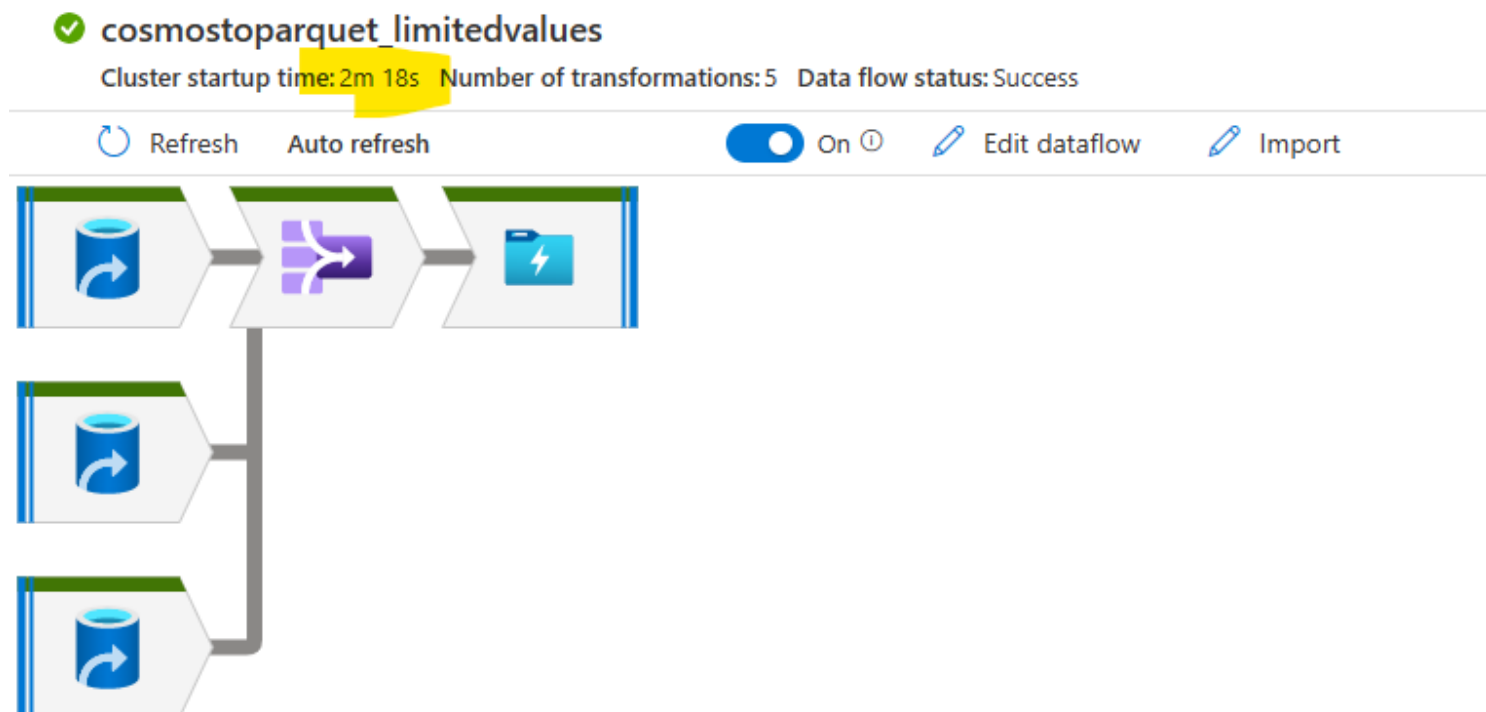
Pipeline run ID: 591592a9-1f24-4282-a16d-b27403a19613

All status ▾

Showing 1 - 3 items

Activity name	Activity type	Run start	Duration	Status	Error	Log	Integration runtime	User properties	Run ID
Advworksdf	Data flow	1/31/2023, 3:39:46 PM	00:00:17	✓ Succeeded			ManualDFIR (West US)		fece985c-7d89-470
dataflow3	Data flow	1/31/2023, 3:39:27 PM	00:00:19	✓ Succeeded			ManualDFIR (West US)		4da2dcc6-08f4-478
cosmostoparquet_limitedvalu...	Data flow	1/31/2023, 3:36:14 PM	00:03:12	✓ Succeeded			ManualDFIR (West US)		64154f29-ba32-49c

The first Dataflow took around 2m 18s to spin the cold cluster as shown below:



✓ **cosmostoparquet_limitedvalues**

Cluster startup time: 2m 18s Number of transformations: 5 Data flow status: Success

Refresh Auto refresh On ⓘ Edit dataflow Import

In this case, although reusing the existing cluster, for second DF activity the dataflow uses warm cluster and it took ~3sec to spin up (shown below). During this time, new containers are being created on that

cluster for isolation, and spark sessions utilize those new isolated containers in the same existing cluster for job processing. ADF - Dataflow keeps all jobs isolated from each other. When we submit a job they create a new container and start spark again for isolation. So you won't find that operation in ADF logs. Customers could use this Scenario to enhance performance.

dataflow3

Cluster startup time: 2s 552ms Number of transformations: 2 Data flow status: Success

Refresh

Auto refresh

On ⓘ

Edit dataflow

Import

We can use the above Kusto query to find the telemetry of detailed timelines of activities.

Scenario: 3

Customers could have a pipeline with multiple Dataflow activities with an Integration Runtime (IR) with specified Time To Live (TTL) and default Integration Runtime set up one for each Dataflow Activity.

For instance, the following Pipeline consists of three Dataflow activities. Out of all activities, 2 DF activities setup with an Azure Integration Runtime with TTL of 10 Minutes while the remaining one activity setup with Default Integration Runtime.

cosmostoparquet_limitedvalues

dataflow3

Advworxsdf

Activity runs

Pipeline run ID 4300374d-88f4-459c-8589-1acb5c73593a

All status

Export to CSV

Showing 1 - 3 items

Activity name	Activity type	Run start	Duration	Status	Error	Log	Integration runtime	User properties	Run ID
Advworxsdf	Data flow	1/31/2023, 3:56:54 PM	00:03:41	Succeeded			AutoResolveIntegrationR		a326546b-a351-49e
dataflow3	Data flow	1/31/2023, 3:56:27 PM	00:00:27	Succeeded			ManualDFIR (West US)		51ca5e80-cad4-456
cosmostoparquet_limitedvalu...	Data flow	1/31/2023, 3:52:56 PM	00:03:31	Succeeded			ManualDFIR (West US)		86380692-9ef9-40d

We can use the above Kusto query to find the telemetry of the completion time.

How to check Cluster Compute Acquisition Duration:

Please refer to this [wiki](#)

How TTL Works?

Please refer to this [wiki](#)

Possible Questions and Recommendations:

1. What exactly happens during Cluster startup time besides cluster creation; warmup? Do we submit the job to the cluster? **Ans: See explanation in Scenario 1 and Scenario 2**
2. Can cluster startup time can be reduced/avoided in Scenario 2? **Ans: No, see explanation in Scenario 2**

Additional Information:

How good have you found this content?

