Determine the recovery time

Last updated by | Keith Elmore | Apr 5, 2021 at 7:56 AM PDT

Contents

Determine the recovery time

How can we determine the recovery time

The recovery time to restore a database by using automated database backups is affected by several factors:

- The size of the database.
- The compute size of the database.
- The number of transaction logs involved.
- The amount of activity that needs to be replayed to recover to the restore point.
- The network bandwidth if the restore is to a different region.
- The number of concurrent restore requests being processed in the target region.

For a large or very active database, the restore might take several hours. If there is a prolonged outage in a region, it's possible that a high number of geo-restore requests will be initiated for disaster recovery. When there are many requests, the recovery time for individual databases can increase. Most database restores complete in less than 12 hours.

For a single subscription, there are limitations on the number of concurrent restore requests. These limitations apply to any combination of point-in-time restores, geo-restores, and restores from long-term retention backup.

To recover by using automated backups, you must be a member of the SQL Server contributor role in the subscription, or be the subscription owner.

Type	Max # of concurrent requests being processed	Max # of concurrent requests being submitted
Single database (per subscription)	10	60
Elastic pool (per pool)	4	200

Public doc reference

https://docs.microsoft.com/en-us/azure/sql-database/sql-database-recovery-using-backups

How good have you found this content?