

errorCode 138 - Internal Server Error (Error update to 4150 refer to Unexpected Failure or Internal Server Error Debug Troubleshooting wiki)

Last updated by | Jackie Huang | Jan 4, 2022 at 12:24 AM PST

Contents

- [Issue](#)
- [PreCheck Conditions:](#)
- [Scenario 1: Cluster has intensive usage](#)
 - [Troubleshooting](#)
 - [Step: 1](#)
 - [Step: 2](#)
 - [Recommendation](#)
- [Scenario 2: Cluster may get OOM issue.](#)
 - [Troubleshooting](#)
 - [Step:1](#)
 - [Step:2](#)
 - [Step:3](#)
 - [Recommendation](#)

Issue

Customer seeing the following error while running the Data flow job:

```
{"ErrorCode":138,"Message":"Internal Server Error:Please try restarting debug session. If problem persists, co
```



PreCheck Conditions:

1. Check the ASC insights if there is an outage on this DF subscription?
2. Due to limitations in ASC insights auto refreshment, it is recommended to manually run the ASC insight
3. If any service insights found, please make sure to cross-check the Outage/regression impact Start and



NOTE: We may get this error in the following scenarios

Scenario 1: Cluster has intensive usage

Troubleshooting

Step: 1

Use below query and find the detailed error message, Integration Runtime:

```
let runId = ''; //Activity Run ID or Pipeline Run ID
let StartTime = datetime(2020-10-01);
let StartTimeminusone = datetime_add('day',-1,StartTime);
cluster('adfcus.kusto.windows.net').database('AzureDataFactory').ActivityRuns
| union cluster('adfneu.kusto.windows.net').database('AzureDataFactory').ActivityRuns
| where activityRunId =~ runId or pipelineRunId =~ runId
| where TIMESTAMP >= StartTimeminusone
| where status != 'Queued' and status != 'InProgress'
| project TIMESTAMP, category, activityType, pipelineRunId, activityRunId, status, failureType
| join kind= leftouter (
cluster('https://azuredmprod.kusto.windows.net').database('AzureDataMovement').CustomLogEvent
| where PreciseTimeStamp >= StartTimeminusone
| where TraceMessage in ("TransferServiceExecutorExecutionState")
| where EventName in ('SsimDeliveryInfoLogEvent','SsimDeliveryWarningLogEvent')
| where Message has "Status\":":2"
| distinct Message,ActivityId,EventName, TraceMessage) on $left.activityRunId==$right.ActivityId
| extend message = extract("<LogProperties><Text>(.*?)</Text></LogProperties>",1,Message)
| extend payload = parse_json(message)
| extend Outputdetails = tostring(payload['Output'])
| extend Error = tostring(payload['Error'])
| distinct TIMESTAMP, category, activityType, pipelineRunId, activityRunId, status, failureType, Error, Output
```

status	failureType	Error	Output
Failed	SystemError	"InnerExecutionError": {"Code": "138", "Message": "Internal Server Error-Please try restarting debug session. If problem persists, contact customer support."}	"clusterId":"1020-182719-fees331","runPageUrl":"https://eastus2.azuredatabricks.net/?o=2848269061254506#job/26938/run/1","effectiveIntegrationRuntime":"DefaultIntegrationRuntime (East US 2)","executionDuration":"423","durationInQueue":"1"

Done (00:01:44): 1 records

PrimaryResult Query Summary QueryStatus Issues: (0)

SELECTED: 1

Name	Path

Error

```
{
  "InnerExecutionError": {
    "Code": "138",
    "Message": "Internal Server Error-Please try restarting debug session. If problem persists, contact customer support.",
    "Details": null
  }
}
```

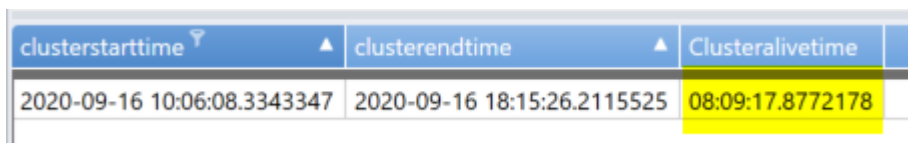
Step: 2

Use below query to find the cluster alive time.

```

let CorrelationId = ''; //Activity Run ID or Pipeline Run ID
let starttime = datetime(2020-09-22);
let endtime = now();
let a = toscalar(cluster('adfcus.kusto.windows.net').database('AzureDataFactory').ActivityRuns
| union cluster('adfneu.kusto.windows.net').database('AzureDataFactory').ActivityRuns
| where PreciseTimeStamp between (starttime..endtime)
//| where category == "SandboxActivityRuns"
| where activityRunId =~ CorrelationId or pipelineRunId =~ CorrelationId
| where status != 'Queued' and status != 'InProgress'
| where activityType == "ExecuteDataFlow"
| join kind=leftouter (
    cluster('azuredmprod.kusto.windows.net').database('AzureDataMovement').TelemetryEvent
    | union cluster('azuredm.kusto.windows.net').database('AzureDataMovement').TelemetryEvent
    | where PreciseTimeStamp between (starttime..endtime)
    | where TelemetryId == "DataflowActivityRunInfo"
    | extend payload = parse_json(Payload)
    | extend runtimeMetrics = payload['RuntimeMetrics']
    | extend ComputeServiceType = payload['ComputeServiceType']
    | where isnotnull(runtimeMetrics)
    | project ActivityId, clusterid = tostring(runtimeMetrics['ClusterId']),ComputeServiceType = tostring(Comp
) on $left.activityRunId == $right.ActivityId
| distinct clusterid, ComputeServiceType,category
| project result = strcat(clusterid,',',ComputeServiceType,',',category));
cluster('adfcus.kusto.windows.net').database('AzureDataFactory').AdfTraceEvent
| union cluster('adfneu.kusto.windows.net').database('AzureDataFactory').AdfTraceEvent
| where env_time between (starttime..endtime)
| extend clusterid = tostring(split(a,',')[0]), computetype=tostring(split(a,',')[1]), category = tostring
| where Message contains clusterid
| summarize clusterstarttime = min(env_time), clusterendtime = max(env_time),Clusterallivetime = max(env_ti

```



clusterstarttime	clusterendtime	Clusterallivetime
2020-09-16 10:06:08.3343347	2020-09-16 18:15:26.2115525	08:09:17.8772178

It seems like this cluster has intensive usage.

Recommendation

Debug sessions are not meant to be used for long periods. Since debug clusters work differently than the job clusters, extended usage could lead to cluster wearing out and eventually run into memory issues and sometimes abrupt restarts. Please consider restarting the debug cluster after some time to free up the old cluster and get a new one.

You may also ask the customer to use activity runs by choosing "Trigger Now" or "Debug" with "Use Activity Runtime" to run their Dataflow, as it would create a new session for every activity run.

Scenario 2: Cluster may get OOM issue.

NOTE: We may not see this OOM error in kusto tables but can see in cluster logs. PG team is working on bubbling up this error in kusto logs. Do not have any ETA at this moment. Reach out to regional EEE's in case if you still need more logs.

Sometimes, the debug run may fail with the same error while the cluster compute unable to handle the data volume/extensive parallel runs. For instance, the below cluster has less alive time but failed with the same error.

Follow below steps for further investigation.

Troubleshooting

Step:1

Check if the cluster alive time is minimal using the query from Step:2 from Scenario 1.

clusterstarttime	clusterendtime	Clusterlvertime
2020-09-24 17:23:47.4098474	2020-09-24 19:02:54.6963468	01:39:07.2864994

Step:2

Check Compute Type and Core Count

```
cluster('adfcus.kusto.windows.net').database('AzureDataFactory').AdfTraceEvent
| union cluster('adfneu.kusto.windows.net').database('AzureDataFactory').AdfTraceEvent
| where TraceCorrelationId == 'XXXXXXXXXXXXXXXXXXXXXXXXXXXX'
| where Message contains "DataFlowActivityPayloadUtilities"
| project env_cloud_location, Message, ComponentId
```

env_cloud_location	Message	ComponentId
eastus2	DataFlowActivityPayloadUtilities.GetDataflowComputePropertiesFromIR[78]: Stop - Returning Compute:DataflowComputeProperties= computeType:General, coreCount:8, nodeCount:0, timeToLive:0, maxConcurrency:..	DataflowManager
eastus2	DataFlowActivityPayloadUtilities.GetComputeLocation[86]: Stop - Returning default compute location: eastus2..	DataflowManager

Step:3

Check How many Activities are running parallelly

In order to find the activities, we need to check first if there is any predecessor pipeline.

```
cluster('adfcus.kusto.windows.net').database('AzureDataFactory').PipelineRuns
| union cluster('adfneu.kusto.windows.net').database('AzureDataFactory').PipelineRuns
| where runId == "XXXXXXXXXXXXXXXXXXXXXXXXXXXX"
| project TIMESTAMP, pipelineName, runId, status, failureType, location, predecessors
```


predecessors			
[{"Name":"377c4e	f0","Id":"377c4ef6-d8	fc0f0","InvokedByType":"PipelineActivity","PipelineName":"PQC	yLd","PipelineRunId":"36f19552-da
[{"Name":"377c4e	f0","Id":"377c4ef6-d8	fc0f0","InvokedByType":"PipelineActivity","PipelineName":"PQC	yLd","PipelineRunId":"36f19552-da
[{"Name":"377c4e	f0","Id":"377c4ef6-d8	fc0f0","InvokedByType":"PipelineActivity","PipelineName":"PQC	yLd","PipelineRunId":"36f19552-da

```
cluster('adfcus.kusto.windows.net').database('AzureDataFactory').ActivityRuns
| union cluster('adfneu.kusto.windows.net').database('AzureDataFactory').ActivityRuns
| where pipelineRunId == "XXXXXXXXXXXXXXXXXXXXXXXXXXXX"
| where status != 'Queued' and status != 'InProgress'
| project TIMESTAMP, pipelineName, pipelineRunId, activityRunId, activityName, activityType, status, category,s
```




timestamp	jobName	pipelineName	activityName	activityType	status	category	start	end	location	correlationId	
2020-09-24 17:48:27.4111978	PO	36f1995	af1e47bf 393c	Trunc ExportCompanyCnt	SqlServerStoredProcedure	Succeeded	SandboxActivityRuns	2020-09-24 17:48:25.8424537	2020-09-24 17:48:27.3138071	eastus2	36f19952-df
2020-09-24 17:48:30.6962269	PO	36f1995	af1e47bf 9708	Lookup Lineage	Lookup	Succeeded	SandboxActivityRuns	2020-09-24 17:48:28.4075711	2020-09-24 17:48:30.5012104	eastus2	36f19952-df
2020-09-24 17:48:31.8633386	PO	36f1995	af1e47bf d28a	Set Lineage	SetVariable	Succeeded	SandboxActivityRuns	2020-09-24 17:48:31.7355875	2020-09-24 17:48:31.7668490	eastus2	36f19952-df
2020-09-24 17:48:34.5575560	PO	36f1995	af1e47bf 56fb	dimItemHierarchy	SqlServerStoredProcedure	Succeeded	SandboxActivityRuns	2020-09-24 17:48:32.4074666	2020-09-24 17:48:34.3567953	eastus2	36f19952-df
2020-09-24 17:56:17.6930841	PO	36f1995	af1e47bf 60a0	dimBusinessUnitId	ExecutePipeline	Succeeded	SandboxActivityRuns	2020-09-24 17:48:32.3449674	2020-09-24 17:56:17.5150332	eastus2	36f19952-df
2020-09-24 17:56:18.1776162	PO	36f1995	af1e47bf c0c0	dimWarehouseId	ExecutePipeline	Succeeded	SandboxActivityRuns	2020-09-24 17:48:32.3449674	2020-09-24 17:56:18.0040067	eastus2	36f19952-df
2020-09-24 17:56:18.3373593	PO	36f1995	af1e47bf 50e0	dimQualityId	ExecutePipeline	Succeeded	SandboxActivityRuns	2020-09-24 17:48:32.3449674	2020-09-24 17:56:18.1520411	eastus2	36f19952-df
2020-09-24 17:56:18.5655049	PO	36f1995	af1e47bf 5737	dimCustomerId	ExecutePipeline	Succeeded	SandboxActivityRuns	2020-09-24 17:48:32.3449674	2020-09-24 17:56:18.3962730	eastus2	36f19952-df
2020-09-24 17:56:18.9963074	PO	36f1995	af1e47bf 282f	dimVendorId	ExecutePipeline	Succeeded	SandboxActivityRuns	2020-09-24 17:48:32.3449674	2020-09-24 17:56:18.7339115	eastus2	36f19952-df
2020-09-24 17:56:19.3766515	PO	36f1995	af1e47bf 36f6	dimSalesOrderId	ExecutePipeline	Succeeded	SandboxActivityRuns	2020-09-24 17:48:32.3449674	2020-09-24 17:56:19.1996340	eastus2	36f19952-df
2020-09-24 17:56:19.7415730	PO	36f1995	af1e47bf 3616	dimAccountId	ExecutePipeline	Succeeded	SandboxActivityRuns	2020-09-24 17:48:32.3449674	2020-09-24 17:56:19.4801762	eastus2	36f19952-df
2020-09-24 17:56:21.3818977	PO	36f1995	af1e47bf 3616	dimAccountId	ExecutePipeline	Succeeded	SandboxActivityRuns	2020-09-24 17:48:32.3449674	2020-09-24 17:56:21.1621154	eastus2	36f19952-df
2020-09-24 17:56:26.4835147	PO	36f1995	af1e47bf efc0	dimInvoiceId	ExecutePipeline	Succeeded	SandboxActivityRuns	2020-09-24 17:48:32.3449674	2020-09-24 17:56:26.2496197	eastus2	36f19952-df
2020-09-24 17:56:32.1721934	PO	36f1995	af1e47bf 9113	dimPurchaseOrderId	ExecutePipeline	Succeeded	SandboxActivityRuns	2020-09-24 17:48:32.3293404	2020-09-24 17:56:31.9689748	eastus2	36f19952-df
2020-09-24 17:56:34.2731272	PO	36f1995	af1e47bf 8ad7	dimEntityId	ExecutePipeline	Succeeded	SandboxActivityRuns	2020-09-24 17:48:32.3449674	2020-09-24 17:56:34.0896796	eastus2	36f19952-df
2020-09-24 17:56:37.6325891	PO	36f1995	af1e47bf 5894	dimDepartmentId	ExecutePipeline	Succeeded	SandboxActivityRuns	2020-09-24 17:48:32.3293404	2020-09-24 17:56:37.4733955	eastus2	36f19952-df
2020-09-24 17:56:40.4543510	PO	36f1995	af1e47bf 3770	dimProductReceiptId	ExecutePipeline	Succeeded	SandboxActivityRuns	2020-09-24 17:48:32.3293404	2020-09-24 17:56:40.2750870	eastus2	36f19952-df
2020-09-24 18:00:43.8062590	PO	36f1995	af1e47bf 1fba	dimResourceId	ExecutePipeline	Succeeded	SandboxActivityRuns	2020-09-24 17:56:18.2031671	2020-09-24 18:00:45.5492754	eastus2	36f19952-df
2020-09-24 18:00:43.3136489	PO	36f1995	af1e47bf a427	dimItemId	ExecutePipeline	Succeeded	SandboxActivityRuns	2020-09-24 17:56:19.9967107	2020-09-24 18:00:48.1195538	eastus2	36f19952-df
2020-09-24 18:00:43.3136489	PO	36f1995	af1e47bf d0e0	dimFinancialTransactionId	ExecutePipeline	Succeeded	SandboxActivityRuns	2020-09-24 17:56:18.7164572	2020-09-24 18:00:48.8651679	eastus2	36f19952-df
2020-09-24 18:01:01.3701717	PO	36f1995	af1e47bf edee	dimSetAsCodeId	ExecutePipeline	Succeeded	SandboxActivityRuns	2020-09-24 17:56:19.9810845	2020-09-24 18:01:01.0583608	eastus2	36f19952-df
2020-09-24 18:01:19.0520992	PO	36f1995	af1e47bf 8c1f	dimInventoryOnHandId	ExecutePipeline	Succeeded	SandboxActivityRuns	2020-09-24 18:00:49.2136247	2020-09-24 18:01:18.7696707	eastus2	36f19952-df
2020-09-24 18:02:00.44432643	PO	36f1995	af1e47bf 7e4c	dimManufacturingTransactionsId	ExecutePipeline	Succeeded	SandboxActivityRuns	2020-09-24 18:00:49.2761240	2020-09-24 18:02:00.1990852	eastus2	36f19952-df
2020-09-24 18:02:32.2914145	PO	36f1995	af1e47bf 23c1	dimProductReceiptId	ExecutePipeline	Succeeded	SandboxActivityRuns	2020-09-24 18:00:49.3386253	2020-09-24 18:02:32.0433732	eastus2	36f19952-df
2020-09-24 18:02:32.8569394	PO	36f1995	af1e47bf 1950	dimQualityId	ExecutePipeline	Succeeded	SandboxActivityRuns	2020-09-24 18:00:49.3386253	2020-09-24 18:02:32.6350693	eastus2	36f19952-df
2020-09-24 18:02:42.2276860	PO	36f1995	af1e47bf 8635	dimItemSpecificationId	ExecutePipeline	Succeeded	SandboxActivityRuns	2020-09-24 18:01:00.1833719	2020-09-24 18:02:41.9892544	eastus2	36f19952-df
2020-09-24 18:02:47.5567545	PO	36f1995	af1e47bf d0a1	dimItemAttributeId	ExecutePipeline	Succeeded	SandboxActivityRuns	2020-09-24 18:00:49.1979982	2020-09-24 18:02:47.315883	eastus2	36f19952-df
2020-09-24 18:03:51.9104477	PO	36f1995	af1e47bf d036	dimSalesId	ExecutePipeline	Succeeded	SandboxActivityRuns	2020-09-24 18:02:43.1371871	2020-09-24 18:03:51.6382866	eastus2	36f19952-df
2020-09-24 18:03:56.4425923	PO	36f1995	af1e47bf 072c	Update Lineage Complete	SqlServerStoredProcedure	Succeeded	SandboxActivityRuns	2020-09-24 18:03:54.2008149	2020-09-24 18:03:56.1776795	eastus2	36f19952-df

Recommendation

If we notice the above screenshot, it could be possible “Execute Pipeline” may consist of multiple Dataflow activities which have been triggered at the same time parallelly. Also, it seems like the above pipeline triggered using an interactive cluster (DataflowAzureDataBricksInteractiveCluster) which is not recommended in this case. We recommend using Trigger run or Activity run with more optimized cores. Please refer to Performance and Tuning [doc](#) 

Additional Information:

- **Icm References:** Scenario1: [Icm link](#) 
- **Author:** rakatuko
- **Keywords:**

How good have you found this content?

