

How to handle increase limit and quota requests

Last updated by | Jarod Zhu | Feb 12, 2023 at 8:48 PM PST

Contents

- [Azure Subscription Service Limits, Quotas, and Constraints](#)
 - [Data Factory limits:](#)
 - [Data Factory v2 Limits](#)
 - [Data Factory v1 Limits](#)
- [Some common scenarios and FAQs](#)
 - [We are reaching a maximum threshold of 5000 for our AD...](#)
 - [Can we increase the Max 40 Activities Per Pipeline limit?](#)
 - [Do we support moving Data Factory across Subscriptions?](#)
 - [Do we support moving Data Factory across Regions?](#)
 - [ARM Template limitations](#)
 - [Resource Name Rules for Data Factory](#)
- [Escalating Limit and Quota increase request to Product Gr...](#)

A lot of times customers would like to increase the quota limit for their subscriptions to allow more resources and operations to be performed than the default limit. In such cases, we ask customers to contact support and the request will come to our queue.



Our team only handles Data Factory related requests. If you come across a case wherein a customer is requesting a limit to be increased for service outside Data Factory, kindly re-route the incident to the appropriate team.

You may refer to the following article to understand different Azure Subscription and service limits, quotas, and constraints.

Azure Subscription Service Limits, Quotas, and Constraints

[Azure subscription and service limits, quotas, and constraints](#) 

Data Factory limits:

- [Data Factory v2 Limits](#) 
- [Data Factory v1 Limits](#) 

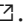
Some common scenarios and FAQs

This section contains some common scenarios and requests which customers have raised with CSS to increase the quotas and limits and response from the Product Group. Please refer to the suggestions and understand the

customer's requirement/roadblock in detail. If there's a hard-stop and customers would still like to increase the limits, please refer below common scenarios:

1. We are reaching a maximum threshold of 5000 for our ADF. Should it be increased?

We have the entity total count limit (default is 5000) to prevent further degradation of the service such as performance on UX. We can increase it but depending on your design you may hit other limits that we cannot change. As a recommendation, we suggest splitting your jobs into 2 different factories; or consider parameterizing pipelines and datasets so you can effectively reuse them. This is not a scalable model for the customer. ADF V2 has a lot of features (parameters, ExecutePipeline, expressions, etc.) and such a large number of entities should not be needed.

You can read about all our limits [here](#) .

The ones that customer should take special attention to if we increase the number of entities are Pipeline concurrency (10K)

- Activity Runs (3K or 1K depending on activity type)
- ARM limits (1200 writes, 12,500 reads per hour)

It is not recommended to go beyond the 5000 entity count limit. If the customer acks to the above information and still want us to increase the max we can do is 10,000 and we won't be able to increase it further.

Please create an AVA thread for this scenario and get this approved from ADFV2 orchestration PG. However, please ask the customer to ACK the above conditions and mention the same in the Ava description before you create an Ava

2. Can we increase the Max 40 Activities Per Pipeline limit?

Most of the times, customers that need more than 40 activities in their pipelines would benefit from refactoring. Take any programming language as an example. You can write one very long method, but it's probably a good idea to identify and refactor help methods. Take a few minutes to see if you are leveraging all ADF features available, like loops, parameters, etc. Smaller and more dedicated pipelines will help you maintaining and troubleshooting your code in the future. It may seem like extra effort initially, but it will definitely be worth it in the long run.

3. Do we support moving Data Factory across Subscriptions?

Yes, Data Factory can be moved across different Resource Groups or Subscriptions. [More Information](#) .


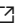
4. Do we support moving Data Factory across Regions?



No, we don't support moving Data Factory across different regions. [More Information](#) .


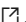
The below table is Work in Progress

Resource	Default limit	Maximum limit	Can limit be increased?	CS
Total number of entities, such as pipelines, data sets, triggers, linked services, Private Endpoints, and integration runtimes, within a data factory	5,000	Contact support  .	Yes, to 10000	Werlin 50 fu de th as or in de yc m lin ca As re we sp jo di fa cc pa pi da ca re is m cu V2 fe (p Ex ex ar nu er nc Th cu

Resource	Default limit	Maximum limit	Can limit be increased?	CS
				ta at inc nu er Pi cc (1 Ru de ac AF wi re lt re gc 5C lin cu th in sti inc we 1C we inc fu cr th sc th fr or PC pl cu AC cc m sa



Resource	Default limit	Maximum limit	Can limit be increased?	CS
				de be cr
Total CPU cores for Azure-SSIS Integration Runtimes under one subscription	64	Contact support  .	Yes	Re Al th wi bu re to re
Concurrent pipeline runs per data factory that's shared among all pipelines in the factory	10,000	10,000	No	
Concurrent External activity runs per subscription per Azure Integration Runtime region  External activities are managed on integration runtime but execute on linked services, including Databricks, stored procedure, Web, and others. This limit does not	3,000	3,000	ask is currently reviewed by PG	

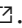
Resource	Default limit	Maximum limit	Can limit be increased?	CS
apply to Self-hosted IR.				
Concurrent Pipeline activity runs per subscription per Azure Integration Runtime region  Pipeline activities execute on integration runtime, including Lookup, GetMetadata, and Delete. This limit does not apply to Self-hosted IR.	1,000	1,000	Yes, to 2000	Re AI ar su AI 
Concurrent authoring operations per subscription per Azure Integration Runtime region  Including test connection, browse folder list and table list, preview data. This limit does not apply to Self-hosted IR.	200	200		
Concurrent Data	Region group 1 ² :	Region group 1 ² : 6,000 Region group 2 ² : 3,000		Re AI

Resource	Default limit	Maximum limit	Can limit be increased?	CS
Integration Units ¹ consumption per subscription per Azure Integration Runtime region 	6,000 Region group 2 ² : 3,000 Region group 3 ² : 1,500 Managed virtual network ² : 2,400	Region group 3 ² : 1,500 Managed virtual network: Contact support  .		
Maximum activities per pipeline, which includes inner activities for containers	40	40	No	M tir th th in wo fro Ta pr lai ex wi lo it': gc id re m fe se le Al av lo pa Sr m pi he m

Resource	Default limit	Maximum limit	Can limit be increased?	CS
				tr yc fu se ef bu de wo lo
Maximum number of linked integration runtimes that can be created against a single self-hosted integration runtime	100	[Contact support] (https://azure.microsoft.com/blog/2014/06/04/azure-limits-quotas-increase-requests/ ).		Re Al th wi bu re to re
Maximum number of node that can be created against a single self-hosted integration runtime	4	Contact support 		
Maximum parameters per pipeline	50	50	No	su Cu br pi di pi


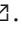
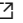
Resource	Default limit	Maximum limit	Can limit be increased?	CS
ForEach items	100,000	100,000	No	
ForEach parallelism	20	50	No	su Cu br pi di pi
Maximum queued runs per pipeline	100	100		
Characters per expression	8,192	8,192		
Minimum tumbling window trigger interval	5 min	15 min		
Maximum timeout for pipeline activity runs	7 days	7 days		
Bytes per object for pipeline objects ³	200 KB	200 KB		
Bytes per object for dataset and linked service objects ³	100 KB	2,000 KB		

Resource	Default limit	Maximum limit	Can limit be increased?	CS
Bytes per payload for each activity run ⁴	896 KB	896 KB		
Data Integration Units ¹ per copy activity run	256	256		
Write API calls	1,200/h	1,200/h This limit is imposed by Azure Resource Manager, not Azure Data Factory.		
Read API calls	12,500/h	12,500/h This limit is imposed by Azure Resource Manager, not Azure Data Factory.		
Monitoring queries per minute	1,000	1,000		
Maximum time of data flow debug session	8 hrs	8 hrs		
Concurrent number of data flows per integration runtime	50	Contact support  .		
Concurrent number of data flows per	20	Contact support  .		

Resource	Default limit	Maximum limit	Can limit be increased?	CS
integration runtime in managed vNet				
Concurrent number of data flow debug sessions per user per factory	3	3		
Data Flow Azure IR TTL limit	4 hrs	4 hrs		
Meta Data Entity Size limit in a factory	2 GB	Contact support  .		

ARM Template limitations

Currently ARM supports only 256 Parameters and Variables in an ARM Template. Moreover, the Template size shouldn't exceed more than 4 MB.

You can exceed some template limits by using a [nested template](#) . For more information, see [Use linked templates when you deploy Azure resources](#) . To reduce the number of parameters, variables, or outputs, you can combine several values into an object. For more information, see [Objects as parameters](#) .

[More Information](#) .

Resource Name Rules for Data Factory

[Resource Name Rules](#) .

Escalating Limit and Quota increase request to Product Group

Kindly reach out to the [Ava](#) for faster responses on the limit quota requests, and contact the respective PGs. Raise an ICM after you've followed this TSG and checking appropriate Product Group in Ava to avoid further delay.

How good have you found this content?

