

Hyperscale FAQ

Last updated by | Pedro Acevedo | Nov 30, 2021 at 9:25 AM PST

1. What happens if storage is full in one worker node?

The node that is full goes to read only. If there is data skew, customer can do rebalance. If there is no data skew and all other workers are nearly full, then Cx needs to upscale storage or add a node and then run rebalance. Other alternatives are to manually move a shard to a node with less usage. This can impact performance however. Also, if colocated tables are move out of place, join queries can fail.

2. How long does it take to restart a server group?

It's not easy to say. If the cluster doesn't have much load during restart, it should be around a minute. Depending on the load (esp. in write-heavy workloads), it can be up to 10 minutes.

3. What rebalance strategies are there?

Command	Is online?	Additional Information
SELECT rebalance_table_shards('distributed_table_name');	Yes	You only need to select one table from the collocation group and rebalance will run for all the tables in that group. This operation won't work if there are no primary keys
SELECT rebalance_table_shards('test_table', shard_transfer_mode => 'force_logical');	No	This is recommended for tables that don't receive DELETE or UPDATE statements. Only READ and INSERT operations will work during this rebalancing strategy
SELECT rebalance_table_shards('test_table', shard_transfer_mode => 'block_writes');	No	This strategy blocks all WRITES and use COPY to replicate the data. READ queries will continue to work
SELECT master_move_shard_placement(12345, 'from_host', 5432, 'to_host', 5432);	No	This function moves a given shard (and shards collocated with it) from one node to another. You can use the following arguments: shard_transfer_mode: (Optional) Specify the method of replication, whether to use PostgreSQL logical replication or a cross-worker COPY command. The possible values are: auto: Require replica identity if logical replication is possible, otherwise use legacy behavior (e.g. for shard repair, PostgreSQL 9.6). This is the default value. force_logical: Use logical replication even if the table doesn't have a replica identity. Any concurrent update/delete statements to the table will fail during replication. block_writes: Use COPY (blocking writes) for tables lacking primary key or replica identity.

Note: If you cancel (CTRL+C) an ongoing rebalance operation, the shards are not cleared in the target node.

4. Is scaling compute an online operation?

No, there is brief downtime.

5. Is scaling storage in the worker nodes or coordinator an online operation?

No, there is brief downtime.


6. How can I scale storage above 2 TB?

Open an ICM

7. Can we do PITR to another region?

No This is not possible. PITR server should be same region. Even Meru to Marlin migration cannot be cross-region. If customers don't see the PITR option, please open an ICM.

8. What regions is PG HS available in?

Check out [this document](#) . If the region is listed, but customer is still unable to deploy in that region, open an ICM.

10. When will geo-redundant backup be available?

It's plan to land between January and July 2022 (tentative, subject to change).