New Static workflow for ADF-LOG ANALYTICS log stream

Last updated by | Jackie Huang | Jan 4, 2022 at 12:24 AM PST

Contents

Problem origin

Log Analytics moved to Kusto implementation which resulted in the below mentioned problems:

Kusto has a column limit of 500 columns which could not support the then implementation of log mapping between log analytics and diagnostics logs, thus dropping the logs

- Since the customer can send logs from multiple products and multiple resources to same LA workspace, they all would append to same AzureDiagnostics table and may result in total number of top-level columns going above 500
- Also, for any json columns, LA(LogAnalytics) would parse through the key value pairs and promote each one of them to a top level column, which also made this problem of hitting 500 columns worse since adf run logs may contain variables and parameters

Symptoms

The logs are present in Azure monitor logs table (Diagnostics Logs) but not in the Log Analytics workspace (delayed and non-transactional log processing possible, success logs may come before in-progress logs in Log Analytics workspace, or they may be completely missing)

Fixes being done

(No extra effort is required for new customers, but existing customers will need to update their config)

- We now have a static workflow for mapping of ADF events (ActivityRuns, PipelineRuns, TriggerRuns – Diagnostics logs work in terms of events) to isolated Log Analytics tables. This means, when logs are pumped from Azure monitor table to Log Analytics Workspace, they will go to a separate isolated ADF event specific table rather than merging with other logs coming to AzureDiagnostics table inside LogAnalytics. Thus, isolating ADF event logs from other resources/product logs.
- This static workflow will also prevent the parsing of json columns, as we are specifically picking out properties, which we consider important to be promoted to top level columns, for rest, the json columns will be treated as strings.

Expectations

- The existing customer can update to new config in place to send ADF Run logs to log analytics. This will send ADF Run logs to ADF event specific tables in their same Log Analytics Workspace instead of the common table (AzureDiagnostics). The preexisting ADF or other logs in the workspace will not be affected, only the new incoming adf logs will be consumed using new static config
- For new customer from 06/21/19 onwards, the default option of config will be updated to use the new static configuration. Resulting in isolating the ADF logs in their event specific tables, and all other logs going to the common table (AzureDiagnostics).

Current Workaround

Change the diagnostics config for ADF resource to point to different loganalytics workspace so that other resources do not affect the ADF logs.

Once the fixes are deployed, below are the steps to be followed to help existing customers switch to new configuration:

1. As shown in below picture, go to the referenced datafactory resource in azure portal, select “diagnostic settings” in left panel, and “edit setting” for the affected Log Analytic workspace (here the datafactory is ADFLALogsTesting and the affected Log Analytic Workspace is mfadfloganalytics)

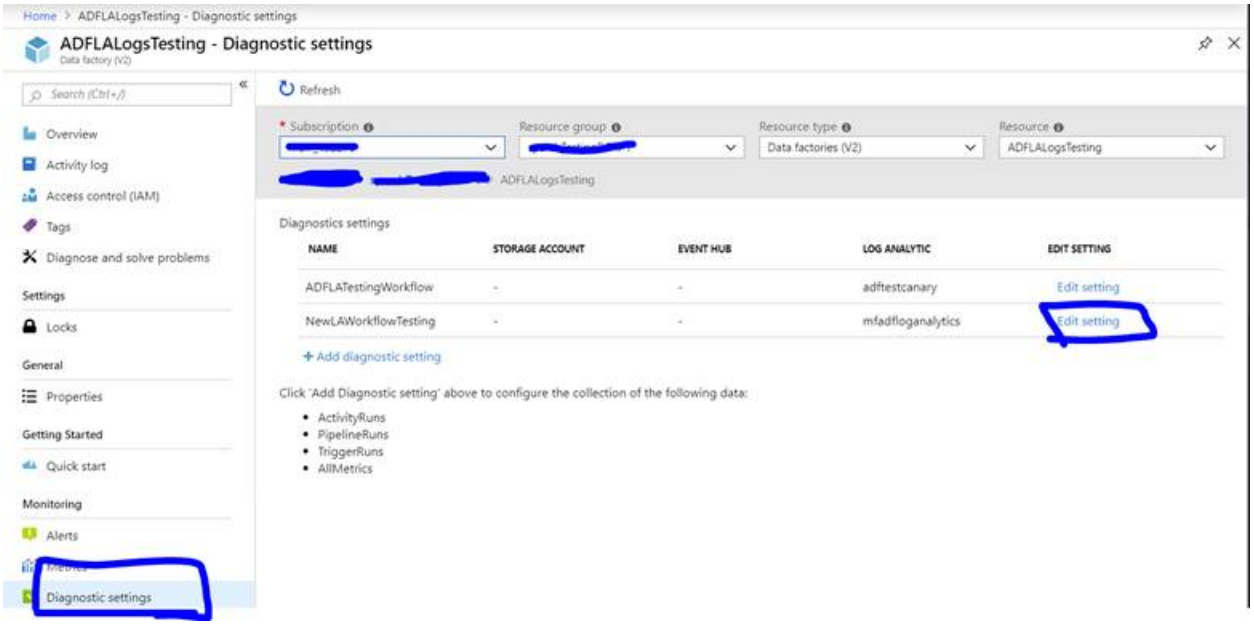

2. Toggle the “destination table” to select the new “Resource Specific” option and save the configuration.

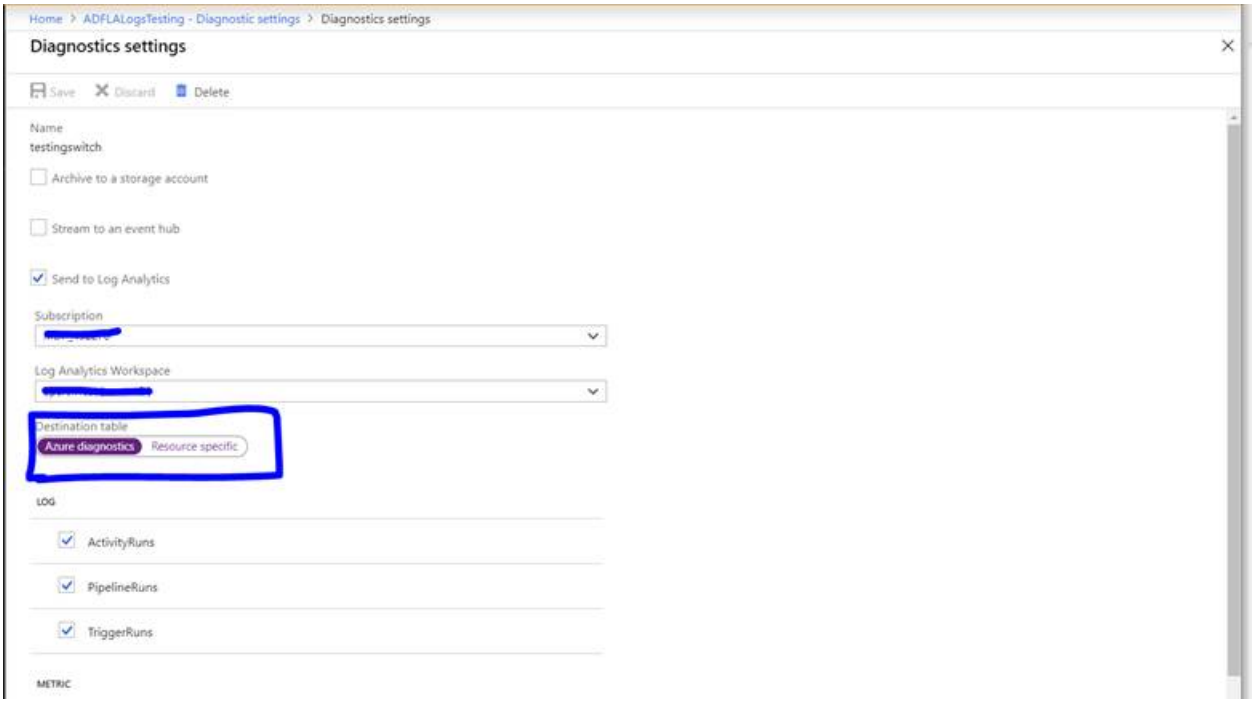

3. From there onwards, the ADF logs will go to ADF event specific tables, “ADFActivityRuns”, “ADFPipelineRuns” and “ADFTriggerRuns” respectively, instead of going to a common table “DiagnosticsLogs”, other resource logs will not be affected.

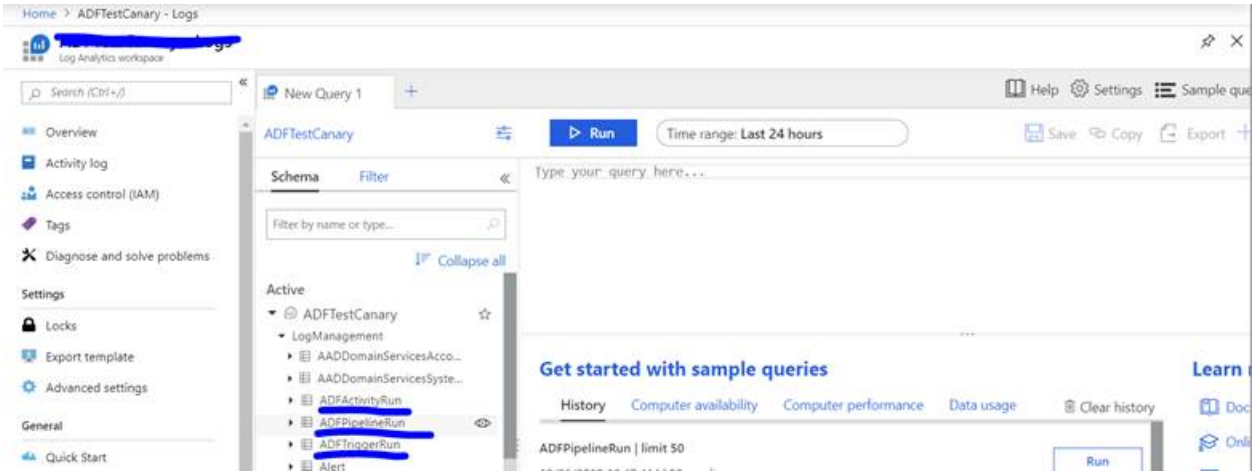

Additional Information:

- **Icm References:**
- **Author:** spagarwa
- **Reviewer:** vimals
- **Keywords:**

How good have you found this content?