

Autovacuum

Last updated by | Dufred Odige | Sep 12, 2021 at 8:38 AM PDT

Autovacuum

Thursday, August 8, 2019
12:06 PM

This is a very helpful explanation for autovacuum
<https://www.2ndquadrant.com/en/blog/autovacuum-tuning-basics/>

We can suggest using table level vacuum settings:

```
autovacuum_vacuum_threshold  
autovacuum_analyze_threshold  
autovacuum_vacuum_cost_delay  
autovacuum_vacuum_cost_limit  
autovacuum_max_workers
```

Best practices:

```
autovacuum_vacuum_cost_delay=10  
autovacuum_vacuum_cost_limit=1000  
autovacuum_vacuum_scale_factor = 0.01  
autovacuum_vacuum_threshold = 1  
autovacuum_analyze_scale_factor = 0.01  
autovacuum_analyze_threshold = 1
```

Auto Vacuum not running/effective

Problem Statement: Why my dead tuples are not cleaned, is Auto Vacuum running at all?

We get this request often from our Customers and I will try to list some of the common causes for such incidents

1. There are too many dead tuples in my z, why is the Auto Vacuum not cleaning them?
2. I see a big bloat in my database, isn't vacuum supposed to clean?
3. My query plans are bad, why isn't vacuum updating stats?
4. Are we seeing XID wraparound warning and transactions failing?

This is not a comprehensive guide to solve all vacuum issues, but a good place to start. Here are the common debugging steps.

1. Is vacuum running? Check if auto vacuum daemon is running at all, default is "ON", connect to the server using SOP0035 How to connect Postgres with azure_superuser if Cert Auth is enabled

Psql-> *SELECT * FROM pg_stat_activity where query like '%vacuum%';*

You should see a record, such as, below row

```
[ RECORD 1 ]
```

<u>datid</u>	
<u>datname</u>	
<u>pid</u>	100
<u>usesysid</u>	
<u>username</u>	
<u>application_name</u>	
<u>client_addr</u>	
<u>client_hostname</u>	
<u>client_port</u>	
<u>backend_start</u>	2019-06-17 17:13:19.313748+00
<u>xact_start</u>	
<u>query_start</u>	
<u>state_change</u>	
<u>wait_event_type</u>	Activity
<u>wait_event</u>	AutoVacuumMain
<u>state</u>	
<u>backend_xid</u>	
<u>backend_xmin</u>	
<u>query</u>	
<u>backend_type</u>	autovacuum launcher

2. If you see the vacuum running in step (1), we can check the progress of auto vacuum

Psql-> *SELECT * from pg_stat_progress_vacuum;*

You should see a record, such as, below row

```
-[ RECORD 1 ]
```

```
-----
```

<u>pid</u>	104701
<u>datid</u>	16419
<u>datname</u>	analytics
<u>relid</u>	1911284001
<u>phase</u>	scanning heap
<u>heap_blks_total</u>	155738368
<u>heap_blks_scanned</u>	75619358
<u>heap_blks_vacuumed</u>	74784601
<u>index_vacuum_count</u>	6
<u>max_dead_tuples</u>	178956970
<u>num_dead_tuples</u>	32346941

Or an improviser for better format

```
SELECT p.pid, Now() - a.xact_start AS duration,
Coalesce(wait_event_type || '.' || wait_event, 'f') AS waiting, CASE WHEN
a.query ~* '^autovacuum.to prevent wraparound' THEN 'wraparound'
```

```

WHEN a.query ~ '^vacuum' THEN 'user' ELSE 'regular' END AS mode,
p.datname AS DATABASE, p.relid :: regclass AS table, p.phase,
Pg_size_pretty(p.heap_blks_total * Current_setting('block_size') :: INT) AS
table_size, Pg_size_pretty(Pg_total_relation_size(relid)) AS total_size,
Pg_size_pretty(p.heap_blks_scanned * Current_setting('block_size') ::
INT) AS scanned, Pg_size_pretty(p.heap_blks_vacuumed *
Current_setting('block_size') :: INT) AS vacuumed, Round(100.0 *
p.heap_blks_scanned / p.heap_blks_total, 1) AS scanned_pct,
Round(100.0 * p.heap_blks_vacuumed / p.heap_blks_total, 1) AS
vacuumed_pct, p.index_vacuum_count, Round(100.0 *
p.num_dead_tuples / p.max_dead_tuples, 1) AS dead_pct FROM
pg_stat_progress_vacuum p join pg_stat_activity a USING (pid) ORDER
BY Now() - a.xact_start DESC;

```

Sample Output

[RECORD 1]	
-----+-----	
pid	104701
duration	03:21:51.330818
waiting	f
mode	regular
database	analytics
table	events
phase	vacuuming indexes
<u>table_size</u>	1188 GB
<u>total_size</u>	1682 GB
<u>scanned</u>	601 GB
<u>vacuumed</u>	571 GB
<u>scanned_pct</u>	50.0
<u>vacuumed_pct</u>	48.0
<u>index_vacuum_count</u>	6
<u>dead_tup_pct</u>	100.0

Check the phase which it's stuck in, it will offer clues where the auto vacuum is spending most of its time and the possible causes.

Check the history of AutoVacuum using

```

Psql-> SELECT relid , schemaname , relname , seq_scan ,
seq_tup_read , last_vacuum , last_autovacuum , last_analyze,
last_autoanalyze FROM pg_stat_user_tables;

```

In some (or most of the customer) scenarios, we have Auto Vacuum running and progressing fine but the table's dead tuples are

Not cleaned up and/or bloat is increasing.

```

Psql-> SELECT * FROM pg_stat_all_tables where relname = '<>';

```

```

-[ RECORD 1 ]
-----+-----
relid          | 992159
schemaname     | public
relname        | contacts__fulltext
seq_scan       | 0
seq_tup_read    | 0
idx_scan       | 649164
idx_tup_fetch  | 442809300
n_tup_ins      | 0
n_tup_upd      | 0
n_tup_del      | 1816244
n_tup_hot_upd  | 0
n_live_tup     | 410668555
n_dead_tup     | 9615572
n_mod_since_analyze | 18
last_vacuum    |
last_autovacuum |
last_analyze   | 2019-07-01 20:10:49.537791+00
last_autoanalyze |
vacuum_count   | 0
autovacuum_count | 0
analyze_count  | 1
autoanalyze_count | 0

```

- You can see live (currently visible) and dead (will not be seen by anyone) tuples count. One of the reasons why vacuum daemon not cleaning up the table could be `autovacuum_vacuum_scale_factor`, which specifies the fraction of the table size when deciding whether to trigger a VACUUM. In this example, `autovacuum_vacuum_scale_factor` is set to 0.05, the engine calculates the scale factor as $9615572 \div (410668555 + 9615572) = 0.0228$ which is smaller than 0.05, and that's why vacuum is not triggered for this table.
- There is also a minimum value to be met for `autovacuum_vacuum_threshold` when deciding whether to trigger a vacuum, this is in addition to the above calculation, this is to prevent excessive vacuuming i.e. if the threshold prevents vacuum for low values, such as 1 row or 10 rows, but if the threshold is set too high, such as 1 Million, vacuum will not trigger until we have a million dead rows.
- The condition for auto vacuum to trigger on table is $\text{dead_tuples} \geq \text{table_size} * \text{scale_factor} + \text{threshold}$. Based on (3) and (4) please tune the config parameters for the vacuum to take effect. In some cases, where an immediate mitigation is needed, you can issue vacuum manually.

Psql-> VACUUM ANALYZE

Scenarios where we have tuples not cleaned up, and still holding up Transaction Id preventing the XID wrap around. This manifests as severe downtime as no new connections are accepted by the engine.

Psql-> SELECT pg_database.datminmxid; <-- Find the database with the smallest value

Psql-to-above-db> SELECT pg_class.relminmxid; <-- To find the culprit table

Check if there is still any open transaction(neither committed nor rolled back) that might be preventing the cleaning.

Psql-> SELECT pid, datname, username, state, backend_xmin

```
FROM    pg_stat_activity
WHERE    backend_xmin IS NOT NULL
ORDER BY age(backend_xmin) DESC;
```

As a mitigation, close the open transaction and also manually unfreeze the xid.

Psql-> VACUUM FREEZE; <-- This should advance the pg_class.relminmxid of the table

Created with Microsoft OneNote 2016.

How good have you found this content?

