

# Vacuum

Last updated by | Lisa Liu | Nov 6, 2020 at 10:35 AM PST

## Why VACUUM?

Whenever rows in a PostgreSQL table are updated or deleted, dead rows are left behind. VACUUM gets rid of them so that the space can be reused. If a table doesn't get vacuumed, it will get bloated, which wastes disk space and slows down sequential table scans (and – to a smaller extent – index scans).

VACUUM also takes care of freezing table rows so to avoid problems when the transaction ID counter wraps around, but that's a different story (we will talk later)

Normally you don't have to take care of all that, because the autovacuum daemon built into PostgreSQL does that for you.

### The problem:

If your tables get bloated, the first thing you check is whether autovacuum has processed them or not:

```
SELECT schemaname, relname, n_live_tup, n_dead_tup, last_autovacuum
FROM pg_stat_all_tables
ORDER BY n_dead_tup
        / (n_live_tup
        * current_setting('autovacuum_vacuum_scale_factor')::float8
        + current_setting('autovacuum_vacuum_threshold')::float8)
DESC
LIMIT 10;
```

If your bloated table does not show up here, `n_dead_tup` is zero and `last_autovacuum` is NULL, you might have a [problem with the statistics collector](#).

If the bloated table is right there on top, but `last_autovacuum` is NULL, you might need to configure autovacuum to be more aggressive so that it gets done with the table.

But sometimes the result will look like this:

schemaname	relname	n_live_tup	n_dead_tup	last_autovacuum
laurenz	vacme	50000	50000	2018-02-22 13:20:16
pg_catalog	pg_attribute	42	165	
pg_catalog	pg_amop	871	162	
pg_catalog	pg_class	9	31	
pg_catalog	pg_type	17	27	
pg_catalog	pg_index	5	15	
pg_catalog	pg_depend	9162	471	
pg_catalog	pg_trigger	0	12	
pg_catalog	pg_proc	183	16	
pg_catalog	pg_shdepend	7	6	

(10 rows)

Why won't VACUUM remove the dead rows?

VACUUM can only remove those row versions (also known as "tuples") that are not needed any more. A tuple is not needed if the transaction ID of the deleting transaction (as stored in the xmax system column) is older than the oldest transaction still active in the PostgreSQL database (or the whole cluster for shared tables).

This value (22300 in the VACUUM output above) is called the "xmin horizon".

There are three things that can hold back this xmin horizon in a PostgreSQL cluster:

Long-running transactions.

Abandoned replication slots.

Orphaned prepared transactions.

### **Long-running transactions.**

You can find those and their xmin value with the following query:

```
SELECT pid, datname, username, state, backend_xmin
FROM pg_stat_activity
WHERE backend_xmin IS NOT NULL
ORDER BY age(backend_xmin) DESC;
```

You can use the pg\_terminate\_backend() function to terminate the database session that is blocking your VACUUM.

### **Abandoned replication slots.**

A replication slot is a data structure that keeps the PostgreSQL server from discarding information that is still needed by a standby server to catch up with the primary.

If replication is delayed or the standby server is down, the replication slot will prevent VACUUM from deleting old rows.

You can find all replication slots and their xmin value with this query:

```
SELECT slot_name, slot_type, database, xmin
FROM pg_replication_slots
ORDER BY age(xmin) DESC;
```

### **Orphaned prepared transactions.**

During two-phase commit, a distributed transaction is first prepared with the PREPARE statement and then committed with the COMMIT PREPARED statement.

Once a transaction has been prepared, it is kept "hanging around" until it is committed or aborted. It even has to survive a server restart! Normally, transactions don't remain in the prepared state for long, but sometimes things go wrong and a prepared transaction has to be removed manually by an administrator.

You can find all prepared transactions and their xmin value with the following query:

```
SELECT gid, prepared, owner, database, transaction AS xmin  
FROM pg_prepared_xacts  
ORDER BY age(transaction) DESC;
```

Use the [ROLLBACK PREPARED](https://www.postgresql.org/docs/11/sql-rollback-prepared.html) SQL statement to remove prepared transactions.

<https://www.postgresql.org/docs/11/sql-rollback-prepared.html>

**How good have you found this content?**

