

PostgreSQL Cache hit rate

Last updated by | Lisa Liu | Nov 6, 2020 at 10:35 AM PST

PostgreSQL Cache hit rate

Monday, January 13, 2020

4:34 PM

The typical rule of thumb for most applications is that only a fraction of its data is regularly accessed. As with many other things data can tend to follow the 80/20 rule with 20% of your data accounting for 80% of the reads and often times its higher than this. Postgres itself actually tracks access patterns of your data and will on its own keep frequently accessed data in cache. Generally you want your database to have a cache hit rate of about 99%.

You can find your cache hit rate with:

```
SELECT
    sum(heap_blks_read) as heap_read,
    sum(heap_blks_hit)  as heap_hit,
    sum(heap_blks_hit) / (sum(heap_blks_hit) + sum(heap_blks_read)) as ratio
FROM
    pg_statio_user_tables;
```

If you find yourself with a ratio significantly lower than 99% then you likely want to consider increasing the number of vCores which will add more memory to increase the value of cache hit.

Created with Microsoft OneNote 2016.

How good have you found this content?

