

PostgreSQL Point in time restore

Last updated by | Lisa Liu | Nov 6, 2020 at 10:34 AM PST

PostgreSQL Point in time restore

Tuesday, February 4, 2020

5:20 PM

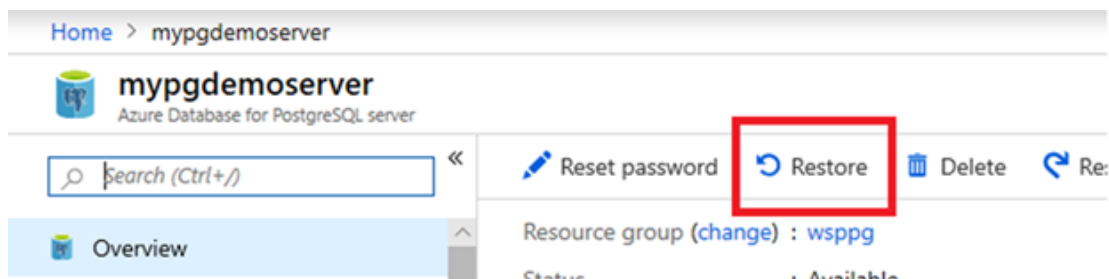
<https://docs.microsoft.com/en-us/azure/postgresql/howto-restore-server-portal>

Point-in-time restore

Independent of your backup redundancy option, you can perform a restore to any point in time within your backup retention period. A new server is created in the same Azure region as the original server. It is created with the original server's configuration for the pricing tier, compute generation, number of vCores, storage size, backup retention period, and backup redundancy option.

Point-in-time restore is useful in multiple scenarios. For example, when a user accidentally deletes data, drops an important table or database, or if an application accidentally overwrites good data with bad data due to an application defect.

You may need to wait for the next transaction log backup to be taken before you can restore to a point in time within the last five minutes.



Restore

In Azure Database for PostgreSQL, performing a restore creates a new server from the original server's backups.

There are two types of restore available:

- **Point-in-time restore** is available with either backup redundancy option and creates a new server in the same region as your original server.
- **Geo-restore** is available only if you configured your server for geo-redundant storage and it allows you to restore your server to a different region.

The estimated time of recovery depends on several factors including the database sizes, the transaction log size, the network bandwidth, and the total number of databases recovering in the same region at the same time. The recovery time is usually less than 12 hours.

Created with Microsoft OneNote 2016.

How good have you found this content?

