

profiling and performance tuning on worker nodes

Last updated by | Shital Modi | Nov 17, 2020 at 8:12 AM PST

profiling and performance tuning on worker nodes

Tuesday, February 12, 2019

10:48 AM

i) **pg_stat_statements** to help you identify the time consuming queries:

[pg_stat_statements](#)

ii) **Monitoring locks:**

https://wiki.postgresql.org/wiki/Lock_Monitoring

<https://www.citusdata.com/blog/2018/02/15/when-postgresql-blocks/>

<https://www.citusdata.com/blog/2018/02/22/seven-tips-for-dealing-with-postgres-locks/>

iii) **Monitoring bloat:**

<https://www.citusdata.com/blog/2017/10/20/monitoring-your-bloat-in-postgres/>

v) **Index usage stats:**

Postgres docs on index maintenance:

https://wiki.postgresql.org/wiki/Index_Maintenance

Identifying unused indexes:

<https://gist.github.com/saicitus/45232c6e1614f3d98da785c28eb0e147>

You could stash those indexes which are not used. It should help in improving resource utilization.

iv) **Others:**

Query to check idx scans vs seq scans on tables:

<https://gist.github.com/saicitus/59738de2b09adcae21d10cd1ef319ae1>

Table/Index cache hit rate:

<http://www.craigkerstiens.com/2012/10/01/understanding-postgres-performance/>

Created with Microsoft OneNote 2016.

How good have you found this content?

