Server side connection pooling

Last updated by | Lisa Liu | Nov 6, 2020 at 10:34 AM PST

Server side connection pooling

Thursday, August 22, 2019
4:26 PM

I see you mentioned you're already using connection pooling for your application. I'm seeing a very high number of connections to your database. Are you currently using connection pooling inside your application, or are you using a server side connection pooler (such as pgbouncer)?

Typically application side pooling opens a bunch of connections which sit idle to the database. Idle applications consume 10MB of memory per connection and can also create contention for establishing new connections and slow down queries.

You can find your idle connections with the following query:

```
SELECT count(*),state FROM pg_stat_activity GROUP BY 2;
Example output:
 count |         state
-------+------------------------------
     7 | active
    32 | idle
    22 | idle in transaction
     4 | idle in transaction(aborted)
(4 rows)
```

You can learn much more about this at https://techcommunity.microsoft.com/t5/Azure-Database-for-PostgreSQL/Connection-handling-best-practice-with-PostgreSQL/ba-p/790883

Essentially if you see a high number of idle connections, it would be recommended to put a Postgres specific connection pooler in place. We recommend leveraging pgbouncer for this.

You can find steps to configure and install pgbouncer at https://techcommunity.microsoft.com/t5/Azure-Database-for-PostgreSQL/Steps-to-install-and-setup-PgBouncer-connection-pooling-proxy/ba-p/730555

Created with Microsoft OneNote 2016.

How good have you found this content?