

Slow perf when loading data into ADLS Gen1 from non binary data source

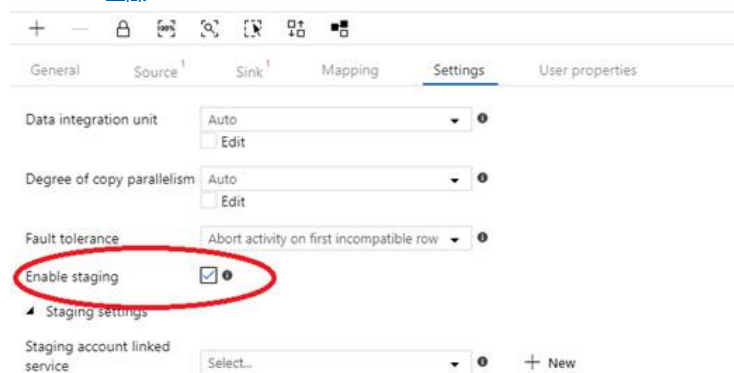
Last updated by | Veena Pachauri | Mar 8, 2023 at 11:57 PM PST

Slow perf when loading data into ADLS Gen1 from non-binary data source

2019年11月13日
10:50 AM

When loading data into ADLS Gen1 from non-binary data source, if the perf bottleneck is in ADLS Gen1 sink side, the customer can try using the Staged Copy to improve the copy performance.

- First, the customer can try this portal to pickup an Azure region that has smallest latency from your local machine: <http://www.azure-speed.com/Azure/Latency>
 - This link can also test the network bandwidth to upload a large file: <http://www.azure-speed.com/Azure/UploadLargeFile>
- Then, the customer can create a Blob Storage account in that region.
- Enable the staged copy and use the new created Blob storage as the staging account. Refer to this doc for more details: <https://docs.microsoft.com/en-us/azure/data-factory/copy-activity-performance#staged-copy>.



The screenshot shows the 'Settings' tab of an Azure Data Factory Copy Activity. The 'Enable staging' checkbox is checked and highlighted with a red circle. Below it, the 'Staging settings' section is expanded, showing a 'Staging account linked service' dropdown menu. Other settings visible include 'Data integration unit' set to 'Auto', 'Degree of copy parallelism' set to 'Auto', and 'Fault tolerance' set to 'Abort activity on first incompatible row'.

Created with Microsoft OneNote 2016.

How good have you found this content?

