

# Configuration (Managed Instance)

Last updated by | Vitor Tomaz | Aug 5, 2020 at 12:42 PM PDT

---

```
//*****
// Configurations, failover, table size
//*****
// C.01
// Primary Node and Secondary Node History and code_package_version history
// B-instance has AppName like b-
// Run this entire with UNION other wise B-instance will mess up the query
// If this query returns no rows, two possibilities
// a. wrong center
// b. standard edition used. in this case, use another query C.01.A
let ResourceStats = materialize(MonDmRealTimeResourceStats
//| where TIMESTAMP >= datetime({StartTime}) and TIMESTAMP <= datetime({EndTime})
| where LogicalServerName =~ "{LogicalServerName}" and database_name =~ "{LogicalDatabaseName}"
| filter replica_type == 0
| summarize AppFirstAppearTime=min(TIMESTAMP), (AppLastAppearTime, physical_database_guid)
=arg_max(TIMESTAMP, physical_database_guid)
    by AppName, LogicalServerName, database_name);
let clPhysicalDbGuid = toscalar(MonDmRealTimeResourceStats | project physical_database_guid);
let dbNameForHadrReplicaStates = iif(isnotempty(clPhysicalDbGuid), clPhysicalDbGuid,
'{LogicalDatabaseName}');
let PoolInfo = ResourceStats
| join kind = inner
(
MonResourcePoolStats
//| where TIMESTAMP > datetime({StartTime}) and TIMESTAMP < datetime({EndTime})
| where LogicalServerName =~ "{LogicalServerName}"
| summarize by LogicalServerName, resource_pool_name, AppName
) on AppName
| join kind= leftouter (
MonAnalyticsRPSnapshot
| where logical_server_name == "{LogicalServerName}"
| summarize min(TIMESTAMP) by logical_server_name,resource_pool_name=name,
resource_pool_dtu_guarantee, database_dtu_cap
) on resource_pool_name
| project AppFirstAppearTime, logical_server_name, database_name, AppName, resource_pool_name,
resource_pool_dtu_guarantee, database_dtu_cap;
(MonDmDbHadrReplicaStates
| where LogicalServerName =~ "{LogicalServerName}" and logical_database_name =~
dbNameForHadrReplicaStates
| where is_primary_replica == 1 and is_forwarder == 0
| where AppName notcontains "b-"
| order by TIMESTAMP asc nulls first
| serialize
| extend prevNodeName=prev(NodeName)
```

```

| extend nodeName = next(NodeName)
| extend isFirst= (NodeName != prevNodeName)
| extend isLast=(NodeName != nodeName)
| where isFirst == true or isLast == true
| extend EndTime=next(TIMESTAMP)
| extend StartTime=TIMESTAMP
| where isFirst ==true
| extend codeversion=extract("^(.*)-WASD", 1, code_package_version)
| extend PrimaryNodeName=NodeName, PrimaryNodeFirstAppearTime=StartTime,
PrimaryNodeLastAppearTime=EndTime
| extend fake="fake"
| project ClusterName, LogicalServerName, logical_database_name, PrimaryNodeName, AppName,
PrimaryNodeFirstAppearTime, PrimaryNodeLastAppearTime, codeversion, fake,
primary_group_database_id=group_database_id
)
| union
(MonDmDbHadrReplicaStates
| where LogicalServerName =~ "{LogicalServerName}" and logical_database_name =~ "
{LogicalDatabaseName}"
| where AppName contains "b-"
| where is_primary_replica == 1 and is_forwarder ==0
| order by TIMESTAMP asc nulls first
| serialize
| extend prevNodeName=prev(NodeName)
| extend nodeName = next(NodeName)
| extend isFirst= (NodeName != prevNodeName)
| extend isLast=(NodeName != nodeName)
| where isFirst == true or isLast == true
| extend EndTime=next(TIMESTAMP)
| extend StartTime=TIMESTAMP
| where isFirst ==true
| extend codeversion=extract("^(.*)-WASD", 1, code_package_version)
| extend PrimaryNodeName=NodeName, PrimaryNodeFirstAppearTime=StartTime,
PrimaryNodeLastAppearTime=EndTime
| extend fake="fake"
| project ClusterName, LogicalServerName, logical_database_name, PrimaryNodeName, AppName,
PrimaryNodeFirstAppearTime, PrimaryNodeLastAppearTime, codeversion,fake,
primary_group_database_id=group_database_id )
| join kind= leftouter (
    MonDmDbHadrReplicaStates
    | where LogicalServerName =~ "{LogicalServerName}" and logical_database_name =~ "
{LogicalDatabaseName}"
    | extend SecondaryNodeName=NodeName
    | where is_primary_replica == 0 and is_forwarder ==0
    | project Secondary_TIMESTAMP=TIMESTAMP, SecondaryNodeName , AppName,
secondary_group_database_id = group_database_id
) on $left.AppName==$right.AppName and $left.primary_group_database_id ==
$right.secondary_group_database_id // group_database_id different of replica is for geo partner
//| extend SecondaryNodeName=case (primary_group_database_id!=secondary_group_database_id and

```

```

PrimaryNodeName==SecondaryNodeName, '', SecondaryNodeName) // Standard edition can have geo
replication
| where (Secondary_TIMESTAMP >= PrimaryNodeFirstAppearTime and Secondary_TIMESTAMP <=
PrimaryNodeLastAppearTime and PrimaryNodeName != SecondaryNodeName) or isnull(SecondaryNodeName)
== true or SecondaryNodeName == ''
| summarize SecondaryNodeNames=makeset(SecondaryNodeName) by bin(PrimaryNodeFirstAppearTime, 1s),
bin(PrimaryNodeLastAppearTime, 1s), ClusterName, LogicalServerName, logical_database_name, AppName,
codeversion, PrimaryNodeName
| join kind= leftouter (
    MonResourcePoolStats
| where LogicalServerName=~ "{LogicalServerName}"
| summarize by AppName, resource_pool_name
) on AppName
| join kind= leftouter (
    PoolInfo
) on AppName
| project ClusterName, AppName, PrimaryNodeName, SecondaryNodeNames,
format_datetime(PrimaryNodeFirstAppearTime, 'yyyy-MM-dd HH:mm:ss'),
format_datetime(PrimaryNodeLastAppearTime, 'yyyy-MM-dd HH:mm:ss'), codeversion, resource_pool_name,
resource_pool_dtu_guarantee, database_dtu_cap
| order by PrimaryNodeFirstAppearTime asc nulls last, PrimaryNodeLastAppearTime asc nulls last

```

```
// C.01.A
```

```
// Primary Node History without cluster
```

```
// use above query when possible
```

```

let ResourceStats = materialize(MonDmRealTimeResourceStats
//| where TIMESTAMP >= datetime({StartTime}) and TIMESTAMP <= datetime({EndTime})
| where LogicalServerName =~ "{LogicalServerName}" and database_name =~ "{LogicalDatabaseName}"
| filter replica_type == 0
| summarize AppFirstAppearTime=min(TIMESTAMP), (AppLastAppearTime, physical_database_guid)
=arg_max(TIMESTAMP, physical_database_guid)
    by AppName, LogicalServerName, database_name);
let clPhysicalDbGuid = toscalar(MonDmRealTimeResourceStats | project physical_database_guid);
let dbNameForHadrReplicaStates = if(isnotempty(clPhysicalDbGuid), clPhysicalDbGuid,
'{LogicalDatabaseName}');
let PrimaryNodes=
(MonDmDbHadrReplicaStates
| where LogicalServerName =~ "{LogicalServerName}" and logical_database_name =~
dbNameForHadrReplicaStates
| where is_primary_replica == 1
| where AppName notcontains "b-"
| order by TIMESTAMP asc nulls first
| serialize
| extend prevNodeName=prev(NodeName)
| extend nextNodeName = next(NodeName)
| extend isFirst= (NodeName != prevNodeName)
| extend isLast=(NodeName != nextNodeName)
| where isFirst == true or isLast == true
| extend EndTime=next(TIMESTAMP)

```

```

| extend StartTime=TIMESTAMP
| where isFirst == true
| extend codeversion=extract("^(.*)-WASD", 1, code_package_version)
| extend PrimaryNodeName=NodeName, PrimaryNodeFirstAppearTime=StartTime,
PrimaryNodeLastAppearTime=EndTime
| project ClusterName, PrimaryNodeName, AppName, PrimaryNodeFirstAppearTime,
PrimaryNodeLastAppearTime, codeversion )
| union
(MonDmDbHadrReplicaStates
| where LogicalServerName =~ "{LogicalServerName}" and logical_database_name =~
dbNameForHadrReplicaStates
| where AppName contains "b-"
| where is_primary_replica == 1
| order by TIMESTAMP asc nulls first
| serialize
| extend prevNodeName=prev(NodeName)
| extend nextNodeName = next(NodeName)
| extend isFirst= (NodeName != prevNodeName)
| extend isLast=(NodeName != nextNodeName)
| where isFirst == true or isLast == true
| extend EndTime=next(TIMESTAMP)
| extend StartTime=TIMESTAMP
| where isFirst == true
| extend codeversion=extract("^(.*)-WASD", 1, code_package_version)
| extend PrimaryNodeName=NodeName, PrimaryNodeFirstAppearTime=StartTime,
PrimaryNodeLastAppearTime=EndTime
| project ClusterName, PrimaryNodeName, AppName, PrimaryNodeFirstAppearTime,
PrimaryNodeLastAppearTime, codeversion );
let AllNodes=
PrimaryNodes
| join kind= leftouter (
    MonDmDbHadrReplicaStates
    | where LogicalServerName =~ "{LogicalServerName}" and logical_database_name =~
dbNameForHadrReplicaStates
    | extend c1=""
    | where is_primary_replica == 0
    | extend SecondaryNodeName=NodeName
    | project TIMESTAMP, SecondaryNodeName , c1 , AppName
) on AppName
| where (TIMESTAMP >= PrimaryNodeFirstAppearTime and TIMESTAMP <= PrimaryNodeLastAppearTime and
PrimaryNodeName != SecondaryNodeName) or isnull(SecondaryNodeName) == true or SecondaryNodeName
== ""
| extend fake="fake"
| summarize SecondaryNodeNames=makeset(SecondaryNodeName) by bin(PrimaryNodeFirstAppearTime, 1s),
bin(PrimaryNodeLastAppearTime,1s), ClusterName, AppName, codeversion, PrimaryNodeName;
AllNodes
| order by PrimaryNodeFirstAppearTime asc nulls last , PrimaryNodeLastAppearTime asc nulls last

```

## // C.02

// Resource Caps such as # of cores, memory, IO etc

MonDmDbResourceGovernance

| where AppName =~ "{AppName}"

| extend MaxCPUCores = case ( primary\_bucket\_fill\_rate\_cpu == 0, (primary\_group\_max\_cpu\*min\_cores)/100 ,  
primary\_bucket\_fill\_rate\_cpu/1000.0)

| extend MaxMemoryMB = case (min\_memory==0, max\_db\_memory/1024.0, min\_memory/1024.0)

| extend MaxMemoryGB = round(MaxMemoryMB /1024.0,2)

| extend MaxIOPS = primary\_group\_max\_io

| extend WorkersLimit = primary\_group\_max\_workers

| extend SessionsLimit=max\_sessions

| extend MaxDbDizeGb = max\_db\_max\_size\_in\_mb/1024

| extend MaxLogKbps = primary\_max\_log\_rate/1024.0

//| distinct slo\_name, MaxCPUCores,MaxMemoryGB, MaxIOPS, MaxLogKbps,MaxDbDizeGb, WorkersLimit,  
SessionsLimit // , LogicalServerName

| summarize StartTime=min(TIMESTAMP), EndTime=max(TIMESTAMP)

by LogicalServerName, AppName, slo\_name, MaxCPUCores,MaxMemoryGB, MaxIOPS,  
MaxLogKbps,MaxDbDizeGb, WorkersLimit, SessionsLimit

| order by StartTime asc

| join kind= leftouter (

MonDmRealTimeResourceStats

| where LogicalServerName =~ "{LogicalServerName}" and database\_name =~ "{LogicalDatabaseName}" and  
AppName =~ "{AppName}"

| filter replica\_type == 0

| summarize arg\_min(TIMESTAMP, \*) by AppName, database\_name

) on AppName

| project StartTime, EndTime, LogicalServerName, AppName, database\_name, slo\_name,  
MaxCPUCores,MaxMemoryGB, MaxIOPS, MaxLogKbps,MaxDbDizeGb, WorkersLimit, SessionsLimit,  
cpu\_cap\_in\_sec

## // C.04

// SLO Change history

MonAnalyticsDBSnapshot

| where logical\_server\_name =~ "{LogicalServerName}" and logical\_database\_name =~ "  
{LogicalDatabaseName}"| summarize min(start\_utc\_date), max(end\_utc\_date) by sql\_instance\_name, AppName, service\_level\_objective ,  
physical\_database\_id, fabric\_partition\_id, logical\_database\_id| project min\_start\_utc\_date, max\_end\_utc\_date, sql\_instance\_name, service\_level\_objective,  
physical\_database\_id,logical\_database\_id,fabric\_partition\_id

| order by min\_start\_utc\_date asc nulls last

## // C.05.A

// Health of replicas

MonDmDbHadrReplicaStates

| where LogicalServerName =~ "{LogicalServerName}" and AppName =~ "{AppName}" //and  
logical\_database\_name =~ "{LogicalDatabaseName}"

| where TIMESTAMP &gt; datetime({StartTime}) and TIMESTAMP &lt; datetime({EndTime})

| where synchronization\_state\_desc == "NOT SYNCHRONIZING" or synchronization\_health\_desc ==

```

"NOT_HEALTHY"
| project TIMESTAMP, NodeName, synchronization_state_desc, synchronization_health_desc
| order by TIMESTAMP asc nulls last
// C.05
// Failover history by WinFab
MonFabricApi
| filter LogicalServerName =~ "{LogicalServerName}" and AppName =~ "{AppName}"
| filter event in ("hadr_fabric_api_replicator_begin_change_role", "hadr_fabric_api_replica_end_change_role")
| where TIMESTAMP > datetime({StartTime}) and TIMESTAMP < datetime({EndTime})
| filter new_role_desc == "PRIMARY"
| project originalEventTimestamp, AppName, event, NodeName, hotpatch_data_package_version,
current_state_desc, new_role, new_role_desc, logical_database_id, physical_database_id, partition_id,
code_package_version, result_desc
| order by originalEventTimestamp asc nulls last

// C.06
// Failover reasons from XTS
// examples: AppMemoryUsageMB, memory usage PLB
WinFabLogs
| where EventType == "Operation" and TaskName == "CRM" | where Id == tolower("{PartitionID}")
| extend decisionId = extract("DecisionId: ([a-z0-9-]+)", 1, Text)
| project decisionId
| summarize count() by decisionId
| join kind = inner (
    WinFabLogs | where EventType == "Decision" and TaskName == "CRM" and TIMESTAMP
> datetime({StartTime}) and TIMESTAMP < datetime({EndTime}) // and ClusterName == "tr6.centralus1-
a.worker.database.windows.net"
    | extend decisionId = Id
    | summarize min(ETWTimestamp), min(Text) by decisionId
) on decisionId
| project PreciseTimeStamp = min_ETWTimestamp, decisionId, Text = min_Text
| order by PreciseTimeStamp asc nulls last
// C.07 MAXDOP
// MAXDOP, 0 means MAXDOP is not set. MAXDOP=# processors
MonDatabaseMetadata
| where LogicalServerName=~ "{LogicalServerName}" and AppName=~ "{AppName}" //logical_db_name =~ "{LogicalDatabaseName}"
| where table_name=="sysobjvalues" and valclass ==7 and objid == 1032 //DB_CONFIG_MAX_DOP
| project TIMESTAMP, MAXDOP=value

// C.08 Database Options
//database options
MonDatabaseMetadata
//| where TIMESTAMP > datetime({StartTime}) and TIMESTAMP < datetime({EndTime})
| where LogicalServerName =~ "{LogicalServerName}" and AppName =~ "{AppName}" and logical_db_name
=~ "{LogicalDatabaseName}"
| where table_name=="sysdbreg"
| extend is_parameterization_forced = (binary_and(status2, 0x08000000) == 0x08000000)

```



```

| extend is_auto_create_stats_on = (binary_and(status2, 0x1000000) == 0x1000000)
| extend is_auto_create_stats_incremental_on = (binary_and(status2, 0x00400000) == 0x00400000)
| extend is_auto_update_stats_async_on = (binary_and(status2, 0x80000000) == 0x80000000)
| extend is_auto_update_stats_async_on = (binary_and(status2, 0x80000000) == 0x80000000)
| extend is_query_store_on = (binary_and(status2, 0x00000010) == 0x00000010) //DBR_QDSENABLED
| extend is_auto_update_stats_on = iff((status2/1073741824) % 2 == 0, false, true)
| project TIMESTAMP, LogicalServerName, logical_db_name, cmptlevel, is_query_store_on,
is_auto_create_stats_on, is_auto_update_stats_on, is_parameterization_forced, is_auto_update_stats_async_on,
is_auto_create_stats_incremental_on
| top 1 by TIMESTAMP desc nulls last

// C.09
//table size
let myAppName="{AppName}";
let PartitionStats=materialize(MonWiDmDbPartitionStats
| where AppName contains myAppName and logical_database_name != 'master' and index_id in (0,1)
| summarize used_page_count=max(used_page_count), row_count=max(row_count) by database_id,
logical_database_name, object_id, index_id, partition_number, data_date=bin(TIMESTAMP, 1d));
let FilteredResults=materialize(MonDatabaseMetadata
| where AppName contains myAppName and logical_db_name != 'master'
| where (table_name=='sysclsobjs' and class==50) or (table_name=='syssschobjs' and ['type']=='U ') or
(table_name=='sysidxstats' and indid in (0,1))
| project TIMESTAMP, table_name, class, ['type'], id, name, nsid, indid);
let schemas=FilteredResults
| where (table_name=='sysclsobjs' and class==50)
| summarize by schema_id=id, schema_name=tolower(name);
let tables=FilteredResults
| where (table_name=='syssschobjs' and ['type']=='U ')
| summarize by schema_id=nsid, object_id=id, table_name=name;
let indexes=FilteredResults
| where (table_name=='sysidxstats' and indid in (0,1))
| extend index_type_desc=iff(['type']==0, 'HEAP', iff(['type']==1, 'CLUSTERED', iff(['type']==5, 'CCI',
tostring(['type']))))
| summarize by object_id=id,index_id=indid,index_type=type,index_type_desc;
tables
| join kind=inner (schemas) on schema_id
| join kind=inner (indexes) on object_id
| join kind=inner (PartitionStats) on object_id
| project database_id, schema_id, object_id, table_type=index_type_desc, partition_number, used_page_count,
row_count, size_kb=used_page_count*8, data_date
//| project database_id, logical_database_name, schema_id, schema_name, object_id, table_name,
table_type=index_type_desc, partition_number, used_page_count, row_count, size_kb=used_page_count*8,
data_date
| summarize argmax(data_date, *) by object_id,partition_number
| top 100 by max_data_date_row_count desc
| project object_id, partition_number, data_date=max_data_date,row_count=max_data_date_row_count,
table_type=max_data_date_table_type

```

```

// C.10
// trace flags
MonNodeTraceETW
| where Message contains 'trace'
| where Message contains "{AppName}"
| where TIMESTAMP > ago(1d)

//*****
// Node wide info
//*****
// N.01
// Perfmon Counter oneminute for CPU and DISK
// TOP N in each category
// watch out
// 1. total CPU %
// 2. individual CPU % pegged to 100% is not good either
MonCounterOneMinute
| where TIMESTAMP > datetime({StartTime}) and TIMESTAMP < datetime({EndTime})
| where NodeName =~ "{NodeName}" and ClusterName =~ "{ClusterName}"
| extend Category = case(CounterName =~ @"\Memory\Available MBytes", "Memory",
    strlen(extract(@"\\Processor\[0-9]{1,2}\\\\% (Privileged Time|Processor Time)", 0, CounterName,
        typeof(string))) > 0, "Processor",
    strlen(extract(@"\\Processor\_Total\\\\% (Processor Time|Privileged Time)", 0, CounterName, typeof(string)))
    > 0, "Total Processor",
    strlen(extract(@"\\LogicalDisk\[A-Z]:\\\\Avg. Disk sec/(Read|Write)", 0, CounterName, typeof(string))) > 0,
    "Disk", "Other")
| where Category != "Other"
| top-nested of bin(TIMESTAMP, 5min) by avg(CounterValue), top-nested of Category by avg(CounterValue),
top-nested 5 of CounterName by Avg_CounterValue=avg(CounterValue) desc
| sort by TIMESTAMP asc nulls last
| project TIMESTAMP, CounterName, Avg_CounterValue
| render timechart

// N.02
// Perfmon Counter oneminute for available MB memory
MonCounterOneMinute
| where TIMESTAMP > datetime({StartTime}) and TIMESTAMP < datetime({EndTime})
| where NodeName =~ "{NodeName}" and ClusterName =~ "{ClusterName}"
| where CounterName =~ @"\Memory\Available MBytes"
| summarize avg(CounterValue) by bin(TIMESTAMP, 5min), CounterName
| render timechart

// N.03
// Node System Event log error summary
MonSystemEventLogErrors
| where TIMESTAMP > datetime({StartTime}) and TIMESTAMP < datetime({EndTime})
| where NodeName =~ "{NodeName}" and ClusterName =~ "{ClusterName}"

```



```
| summarize Count=count(), min(TIMESTAMP), max(TIMESTAMP), min(EventDescription)by EventID  
| order by Count desc nulls last
```

```
// N.04
```

```
// Node Application log error summary
```

```
MonAppEventLogErrors
```

```
| where TIMESTAMP > datetime({StartTime}) and TIMESTAMP < datetime({EndTime})
```

```
| where NodeName =~ "{NodeName}" and ClusterName =~ "{ClusterName}"
```

```
| summarize Count=count(), min(TIMESTAMP), max(TIMESTAMP), min(EventDescription)by EventID
```

```
| order by Count desc nulls last
```

**How good have you found this content?**



-