

Hyperscale - performance troubleshooting diagnostics

Last updated by | Vitor Tomaz | Feb 24, 2023 at 3:28 AM PST

Contents

- [Issue](#)
- [Investigation / Analysis](#)
- [Public Doc Reference](#)

Hyperscale performance troubleshooting guidance


Issue

To troubleshoot performance problems in a Hyperscale database, general performance tuning methodologies on the Azure SQL Database compute node is the starting point of a performance investigation. However, given the distributed architecture of Hyperscale, additional diagnostics have been added to assist. This article points you to Hyperscale-specific diagnostic data.

Investigation / Analysis

As the first steps, run the troubleshooter in ASC and review the report output. It has a separate section for Hyperscale and might already provide you with sufficient information.

Otherwise, follow the general performance troubleshooting methodology based on the performance telemetry provided by the troubleshooter. Check the other performance-related TSGs on this wiki for further steps.

For additional Hyperscale-specific diagnostic steps and data, also go through the public article [SQL Hyperscale performance troubleshooting diagnostics](#)  which has detailed recommendations and background information. It discusses topics like:

- Log rate throttling waits
- Page server reads
- Virtual file stats and IO accounting
- Local RBPEX cache usage
- Data reads and writes
- Log writes
- Data IO in resource utilization statistics

Public Doc Reference

- [SQL Hyperscale performance troubleshooting diagnostics](#) 

How good have you found this content?



-